

Managing Physical Memory

Advanced Operating Systems

Overview

- Physical Memory (PM)
- PM Management in Linux
 - Zones, Nodes, Pages
 - The Memblock Allocator
 - The Buddy Allocator
- PM Management in OpenLSD
 - The Boot Allocator
 - Frame Allocation
- Advanced Linux Features
 - Overlay allocators (slab, vmalloc)
 - Sanity checks

What is “Physical Memory” Anyway?


- **DRAM (Dynamic Random Access Memory)**
 - Organized in memory cells with 1-bit data
 - Each cell has 1 capacitor + 1 transistor
 - Charge/discharge capacitor = 1/0 bit value
 - Capacitors leak charge → Need for periodic refresh
 - Key difference w.r.t. SRAM (used in CPU caches)
 - Typical refresh rates: 8ms - 64ms



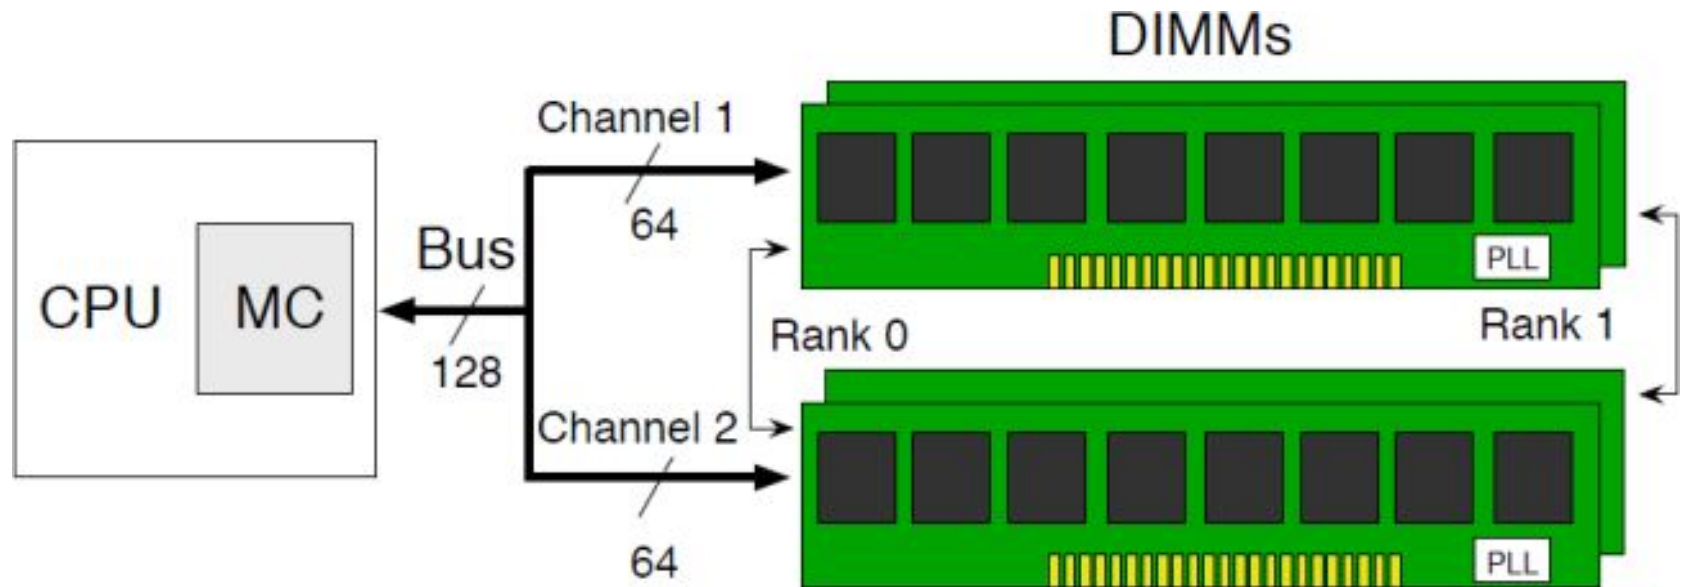
What is “Physical Memory” Anyway?

- **DRAM organization**

- Channel
- DIMM
- Rank
- Chip
- Bank
- Row/Column

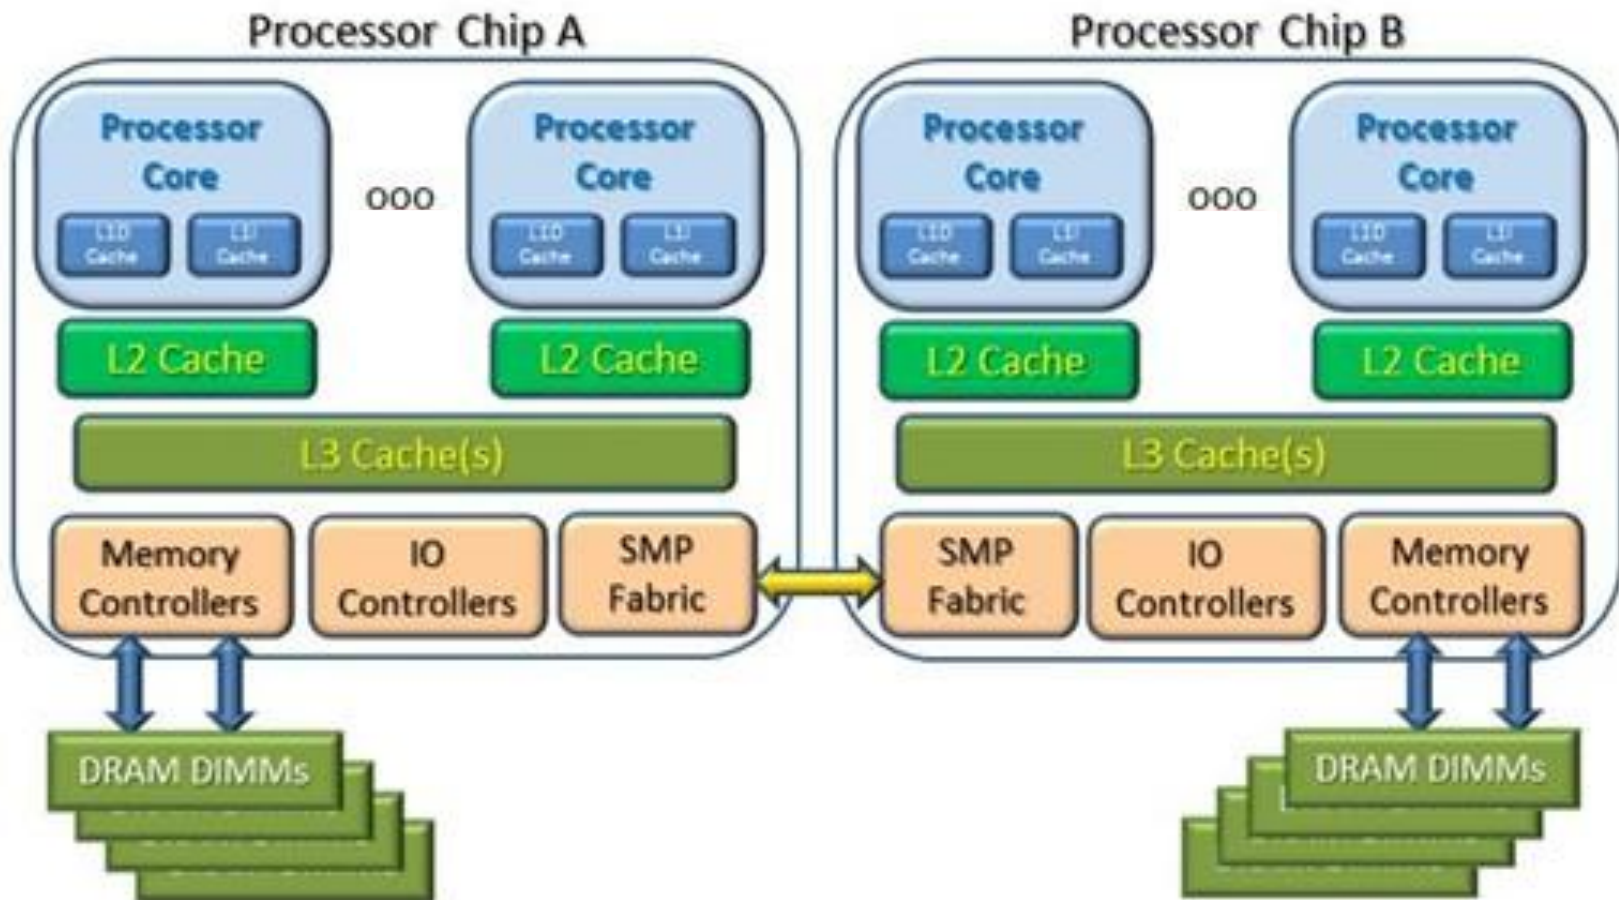


What is “Physical Memory” Anyway?



- 1 / 2 channels, N DIMMs per channel
- 1 / 2 ranks per DIMM

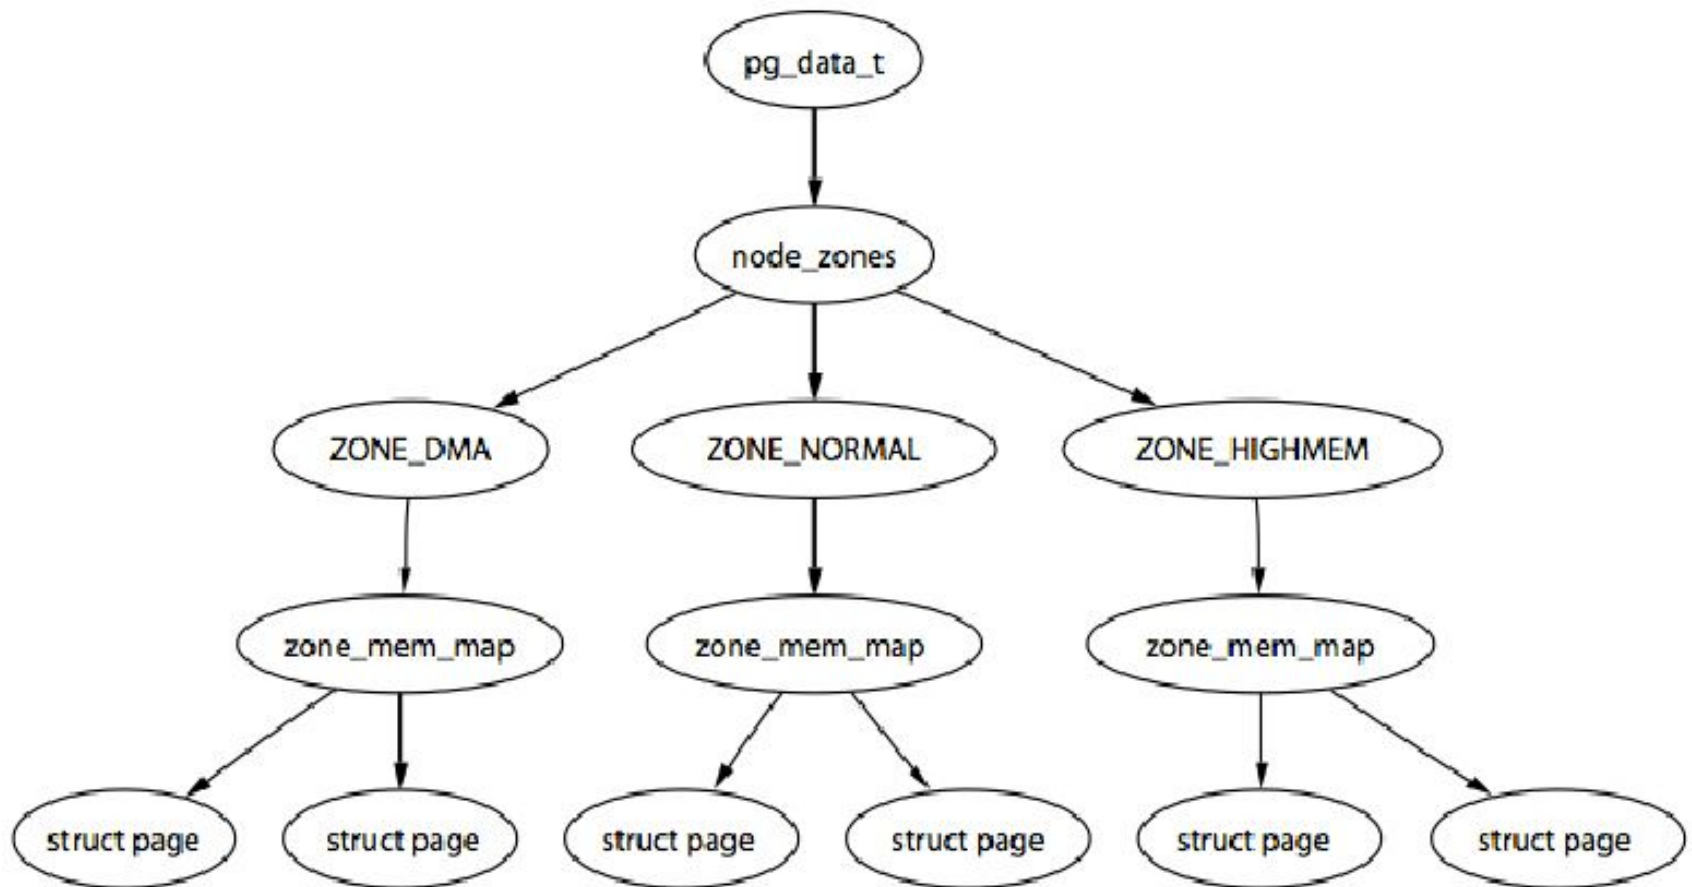
The Big Picture: UMA vs. NUMA



Managing Physical Memory: Linux

- Logically divided in a number of consecutive physical memory pages (page frames) identified by page frame numbers (PFNs)
- How are frames organized on Linux?
 - Nodes
 - The NUMA abstraction of N “banks”
 - Zones
 - Tagged regions for each node
 - Pages
 - Physical page frames for each zone

Managing Physical Memory: Linux



Managing Physical Memory: Nodes

linux/mmzone.h:

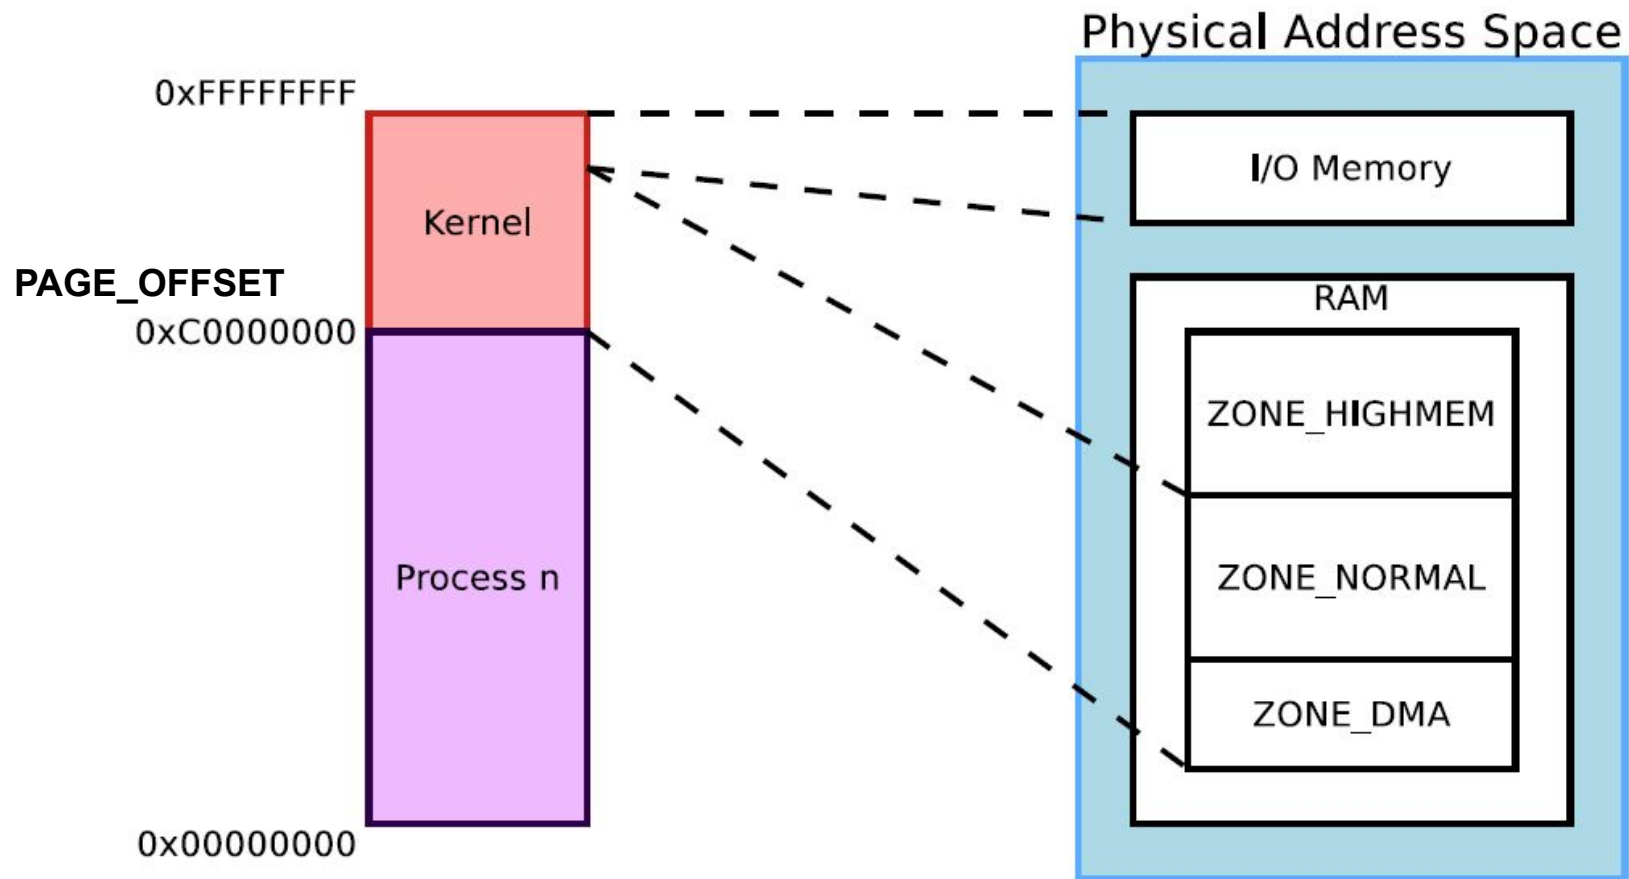
```
662 typedef struct pglist_data {
    struct zone node_zones[MAX_NR_ZONES];
    struct zonelist node_zonelists[MAX_ZONELISTS];
    int nr_zones;
    struct page *node_mem_map;
    unsigned long node_start_pfn;
    unsigned long node_present_pages; /* no. of phys pages */
    unsigned long node_spanned_pages; /* size of physical page
                                        range, with holes */

    int node_id;
#ifdef CONFIG_COMPACTION
    /* ... */
#endif
} pg_data_t;
```

Managing Physical Memory: Zones

- Zone ranges:
 - ZONE_DMA: 0 MB - 16 MB (DMA)
 - ZONE_DMA32: 16 MB - 4 GB (DMA32)
 - ZONE_NORMAL: 16 MB - 896 MB (kernel)
 - ZONE_HIGHMEM: 896 MB - End (kmap, x86)
- Managed independently
- When under memory pressure:
 - Zone boundaries become “blurry”
 - Watermarking strategy to free pages (**kswapd**):
 - WMARK_MIN: Trigger direct reclaim
 - WMARK_LOW: Trigger async reclaim
 - WMARK_HIGH: Stop reclaiming

Managing Physical Memory: Zones (x86)



Managing Physical Memory: Zones

linux/mmzone.h:

```
385 struct zone {
    unsigned long _watermark[NR_WMARK];
    struct pglist_data      *zone_pgdat;
    struct per_cpu_pageset __percpu *pageset;
    unsigned long           zone_start_pfn;
    unsigned long           spanned_pages;
    unsigned long           present_pages;
    const char              *name;
    struct free_area         free_area[MAX_ORDER];
    spinlock_t              lock;
    ZONE_PADDING(_pad2_)
    atomic_long_t           vm_stat[NR_VM_ZONE_STAT_ITEMS];
} ____cacheline_internodealigned_in_smp;
```

Managing Physical Memory: Pages

- 4KB units of physical memory
- struct page objects in a mem_map array

linux/mmzone.h:

```
642      /*      The      array      of      struct      pages      */  
      extern struct page *mem_map;
```

asm/page.h

```
59 #define __va(x) ((void *) ((unsigned long) (x) + PAGE_OFFSET))  
69 #define virt_to_page(kaddr) pfn_to_page(__pa(kaddr) >> PAGE_SHIFT)  
71 extern bool __virt_addr_valid(unsigned long kaddr);
```

asm-generic/memory_model.h:

```
33 #define pfn_to_page(pfn) (mem_map + ((pfn) - ARCH_PFN_OFFSET) / 3)
```

Managing Physical Memory: Pages

linux/mm_types.h:

```

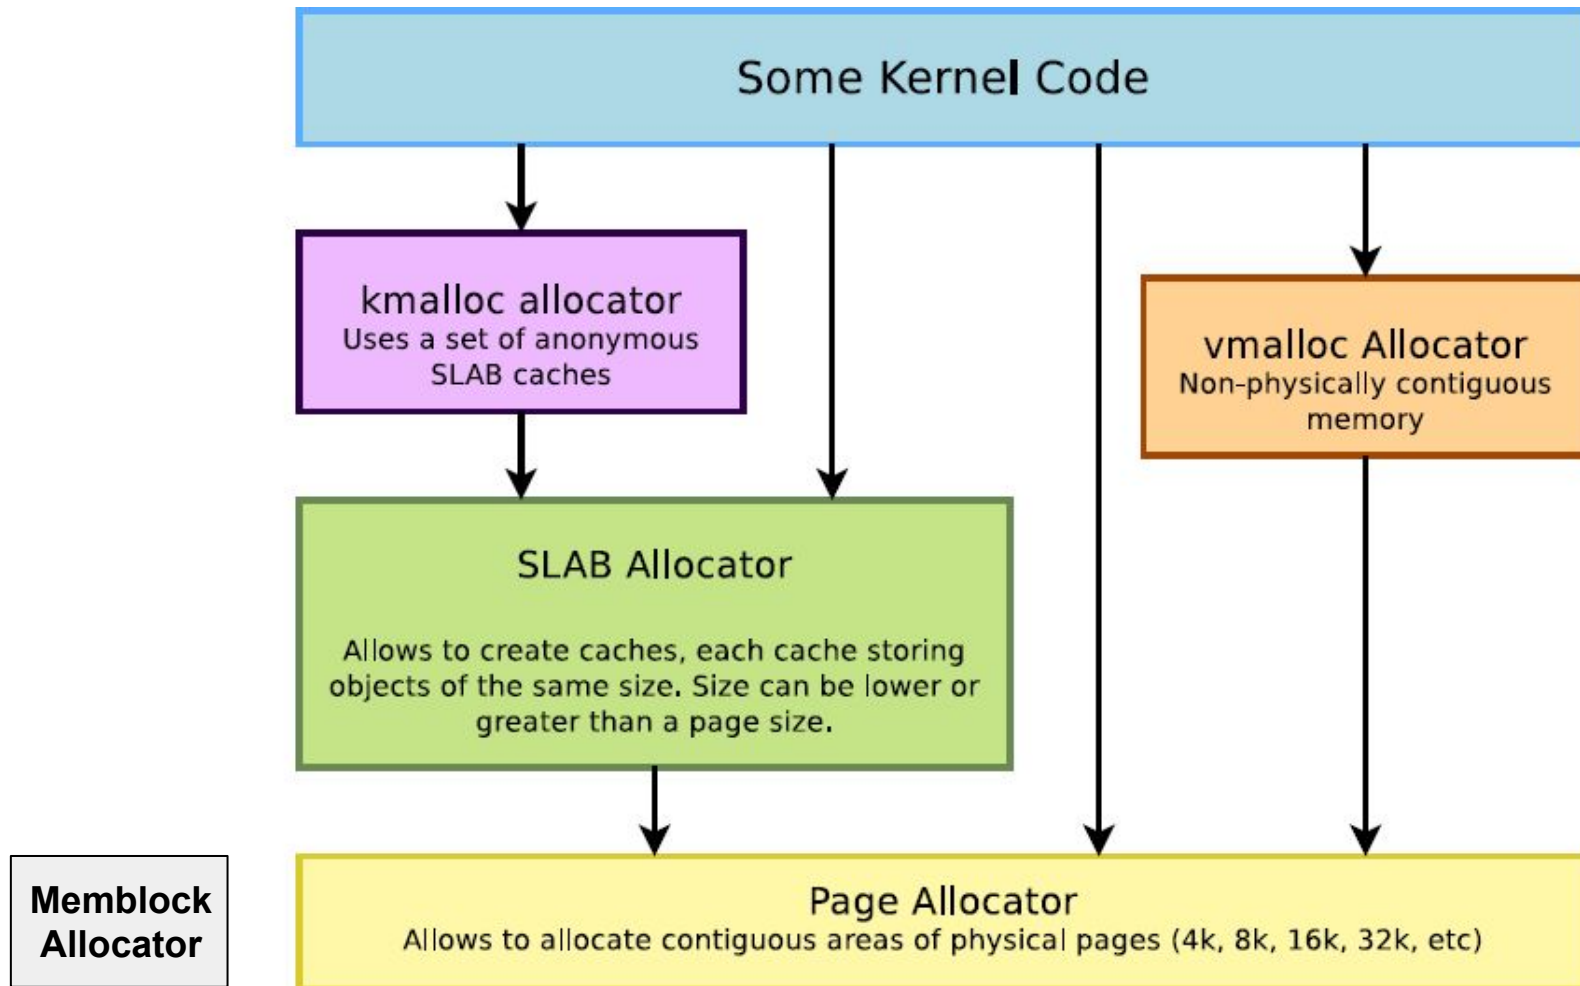
        68                struct                page                {
                unsigned long flags;

                                struct                list_head    lru;
                                struct                address_space  *mapping;
                                                pgoff_t            index;
                                                atomic_t            _refcount;
                                struct                mem_cgroup    *mem_cgroup;
                                                void                *virtual;

        } _struct_page_alignment;
```

\$ cat /proc/*info* | more

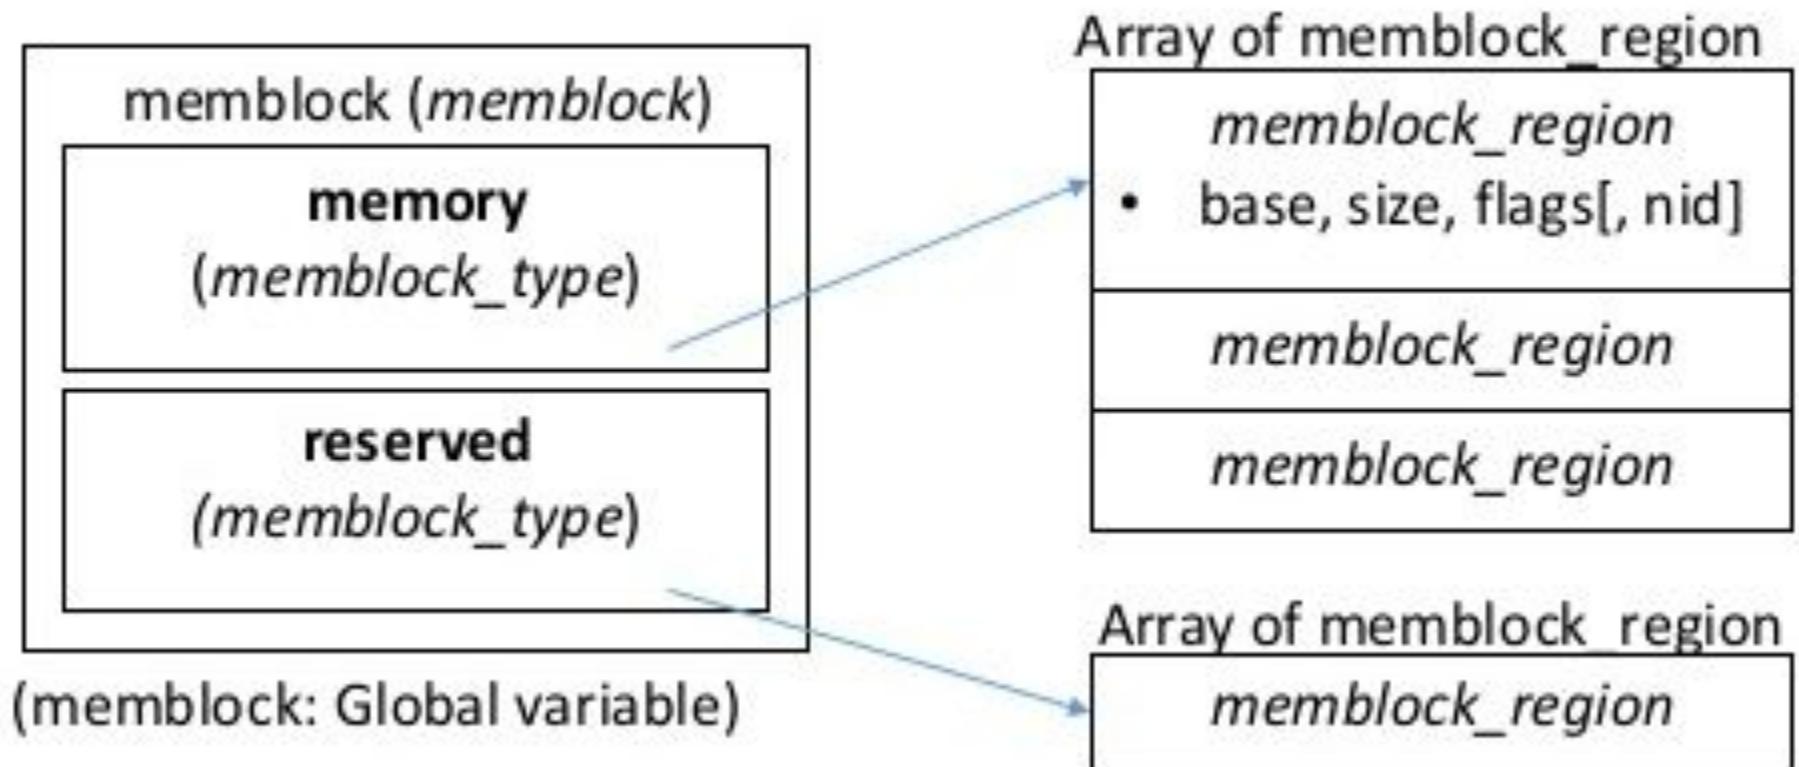
Allocating Memory: Linux Overview



The Linux Memblock Allocator

- Early boot-time (low memory) allocator
- Replaces the old bootmem allocator
- Mostly used to initialize buddy allocator
- Discarded after initialization
- Consists of two main arrays:
 - **memory**: all present memory in the system
 - **reserved**: allocated memory ranges
- Allocates by finding regions in:
 - **memory && !reserved**

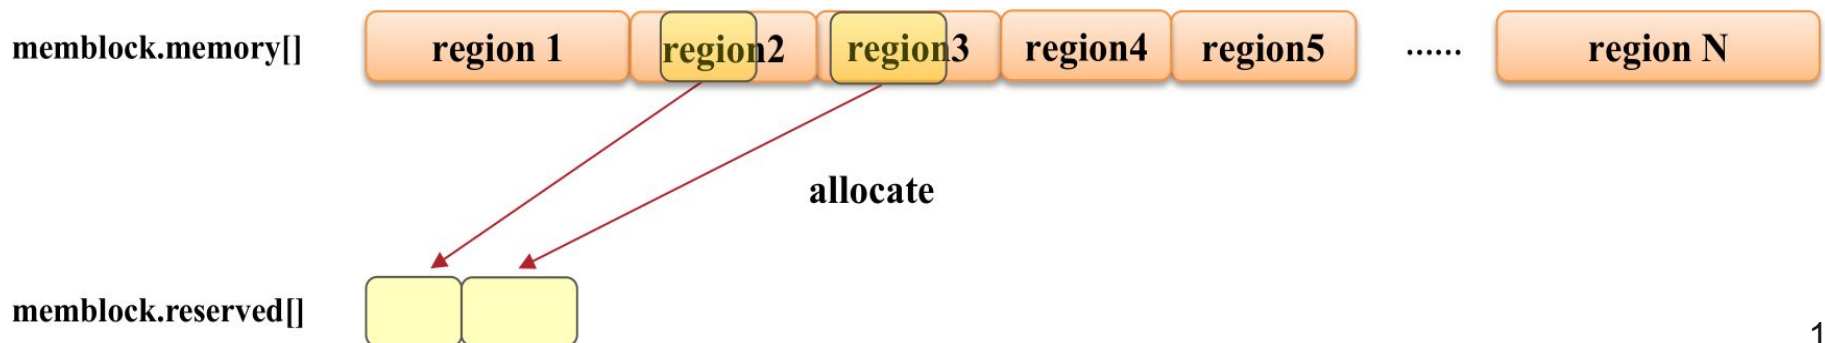
The Linux Memblock Allocator



The Linux Memblock Allocator

linux/memblock.h:

```
49 struct memblock_region {      66 struct memblock_type {
    phys_addr_t base;           unsigned long cnt; /* #regions */
    phys_addr_t size;           unsigned long max; /* alloc size */
    enum memblock_flags flags;  phys_addr_t total_size; /* tot size */
    /* ... */                  struct memblock_region *regions;
};                               char *name;
                                };
```

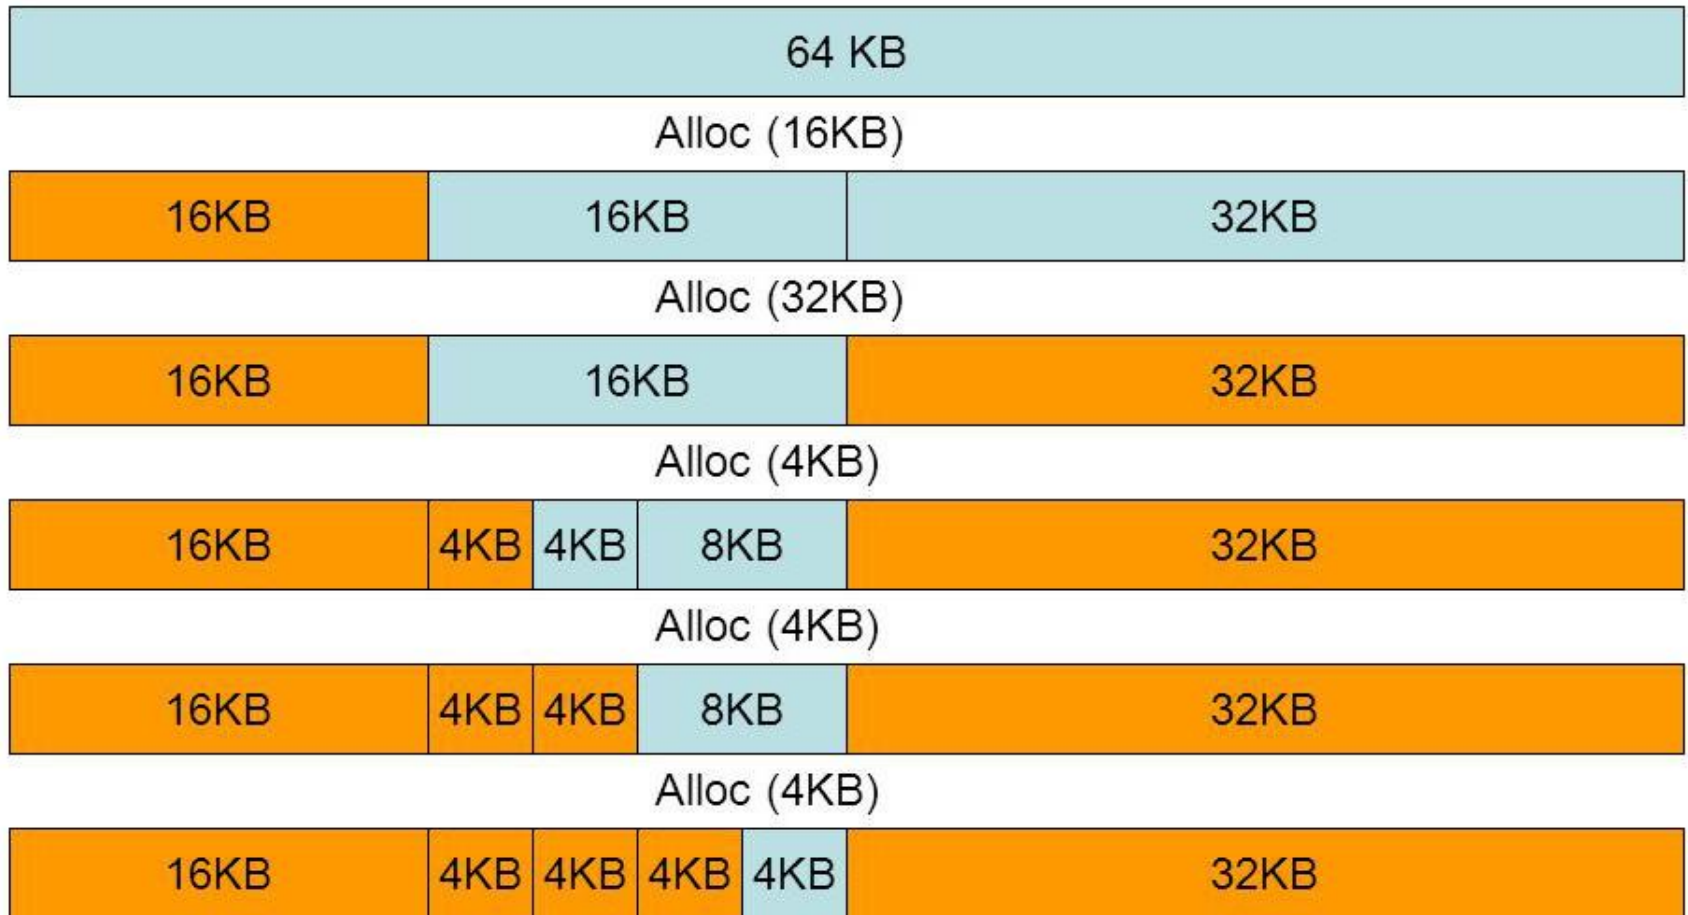


The Linux Memblock Allocator

- Setup:
 - Add all available phys memory regions to memory
 - Add “reserved” ones to reserved
 - All regions sorted by base address
- Allocation:
 - First-fit memory && !reserved
 - Simply memblock_add_range to reserved
 - Merge neighboring regions as necessary
- Deallocation:
 - Linear scan in reserved for containing region
 - Simply memblock_remove_range from reserved
 - First split region if necessary

The Linux Buddy Allocator

- Power-of-two allocator with free coalescing
- Blocks are arranged in 2^N (N =order) pages
- Allocations are satisfied by exact N
 - If not possible, split larger $2^{(N+1)}$ block to 2×2^N
 - The two smaller blocks are called *buddies*
 - If not possible, split larger $2^{(N+2)}$ block twice, etc...
- Deallocations return block to allocator
 - If buddy is free, coalesce into a larger block
 - If larger block's buddy is free, coalesce again, etc...
- Behavior in short: recursive 2^N split/merge operations on allocation/deallocation



The Linux Buddy Allocator: Split (and Merge)

The Linux Buddy Allocator: Internals

- Per node,zone array of MAX_ORDER freelists
 - Freelist N maintains free blocks of size 2^N
 - Split/merge: moves block(s) to previous/next freelist

`linux/mmzone.h:`

```
98 struct free_area {  
    struct list_head free_list[MIGRATE_TYPES];  
    unsigned long    nr_free;  
};
```

- PG_buddy, order in page attributes (set if free)
- Buddy operations
 - $O(1)$ *reserve* (update page flags)
 - $O(1)$ *lookup* (flip “order” bit in block address)

Order

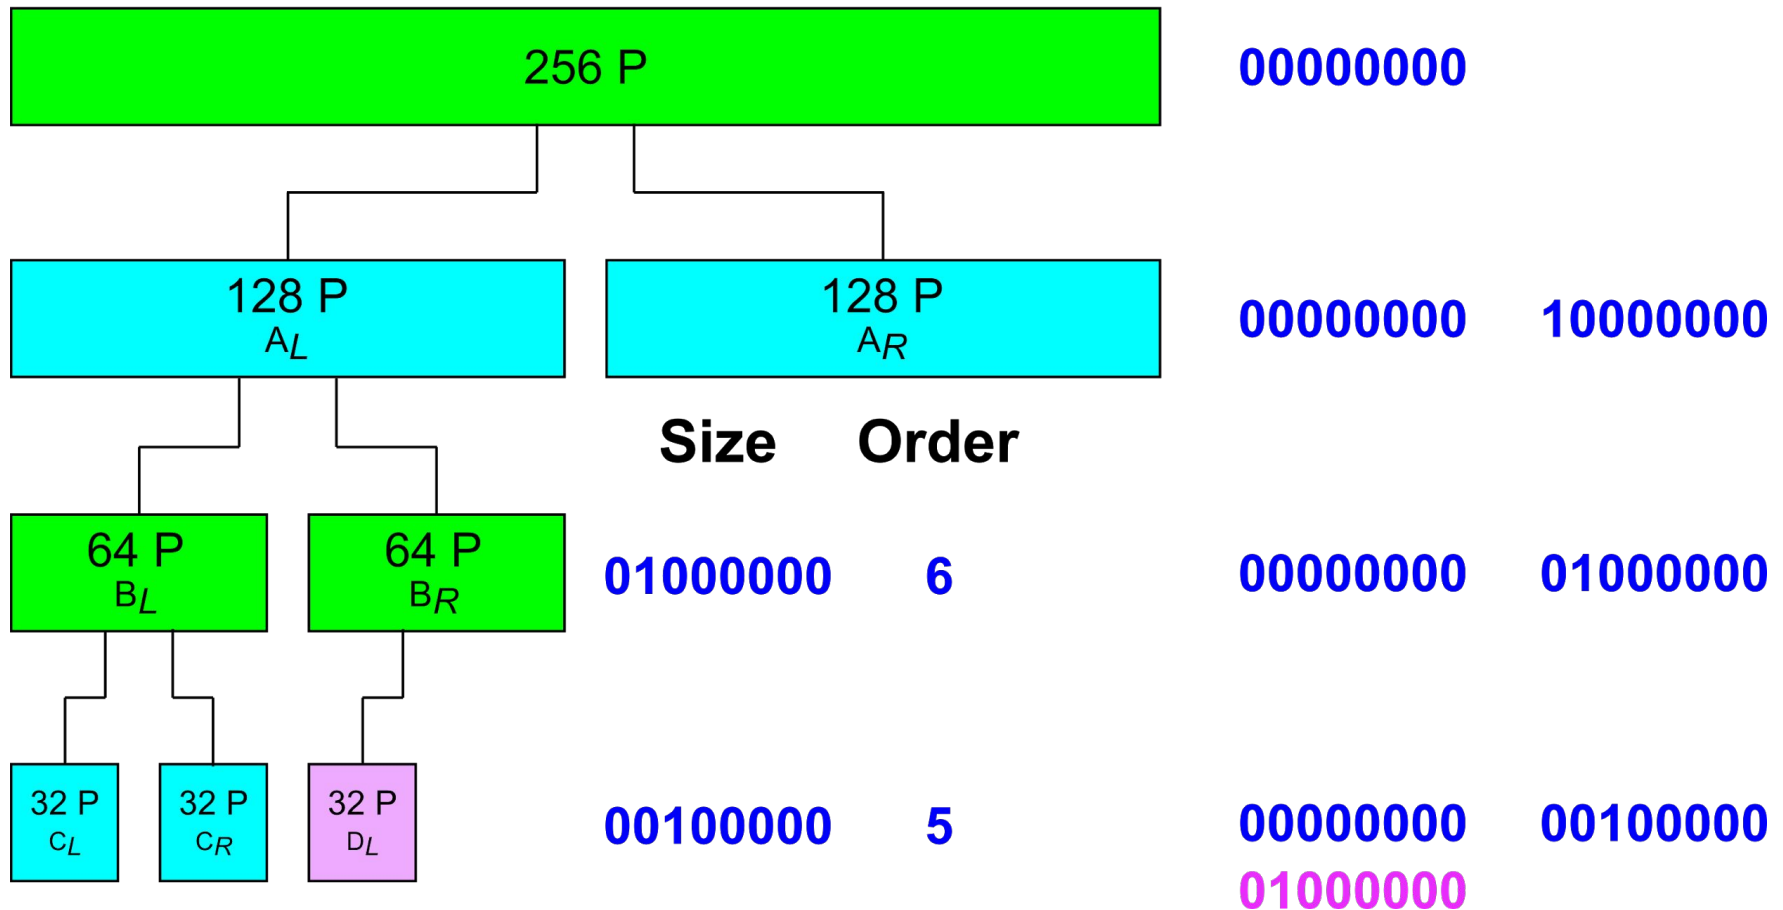
2^{10}	1024 Pages
2^9	512Pages
2^8	256Pages
2^7	128Pages
2^6	64Pages
2^5	32Pages
2^4	16Pages
2^3	8Pages
2^2	4 Pages
2^1	2 Pages
2^0	Page

```
#define MAX_ORDER 11
/* Default setting */
```

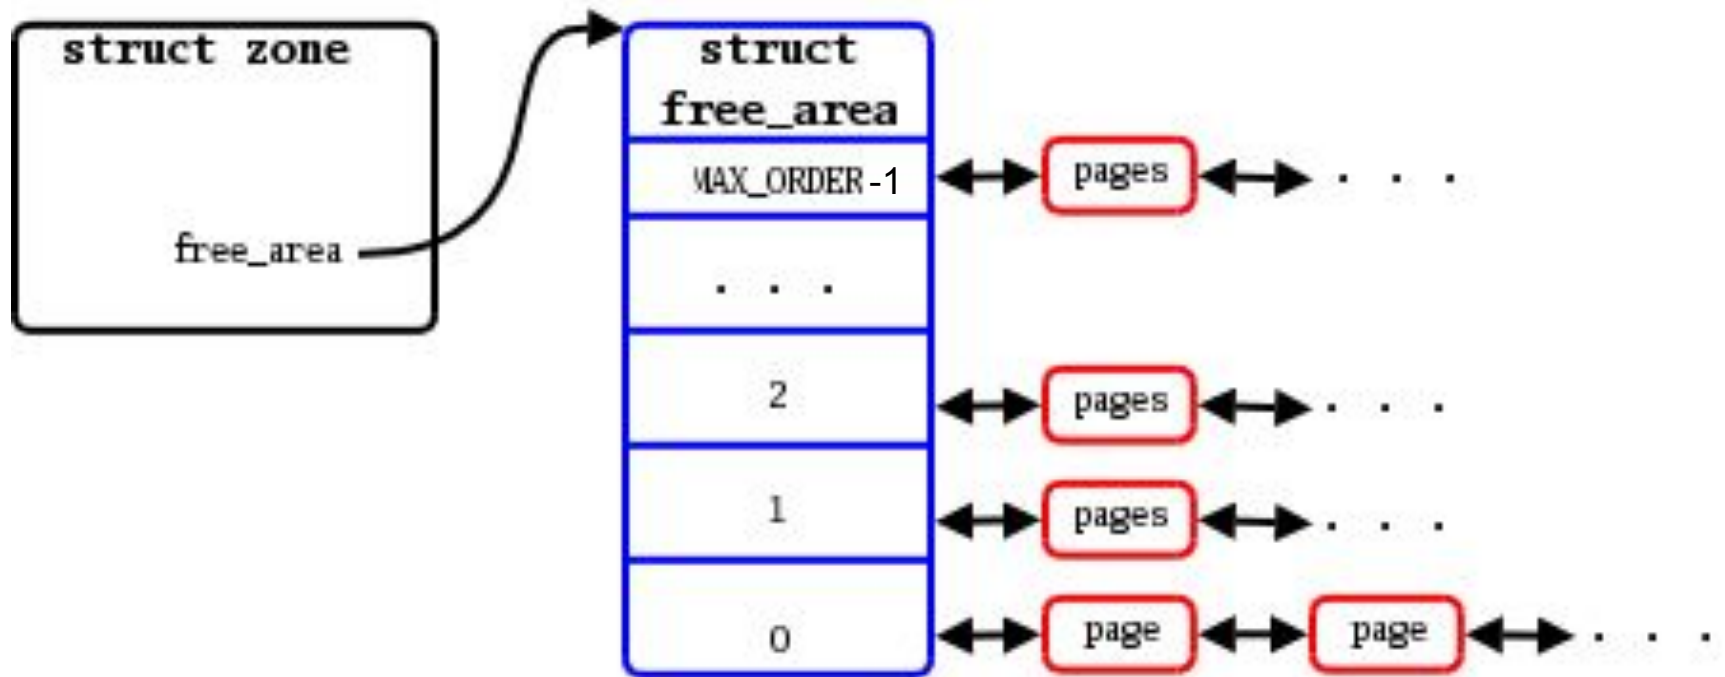
The Linux Buddy Allocator: Per-order Freelists

Physically contiguous pages

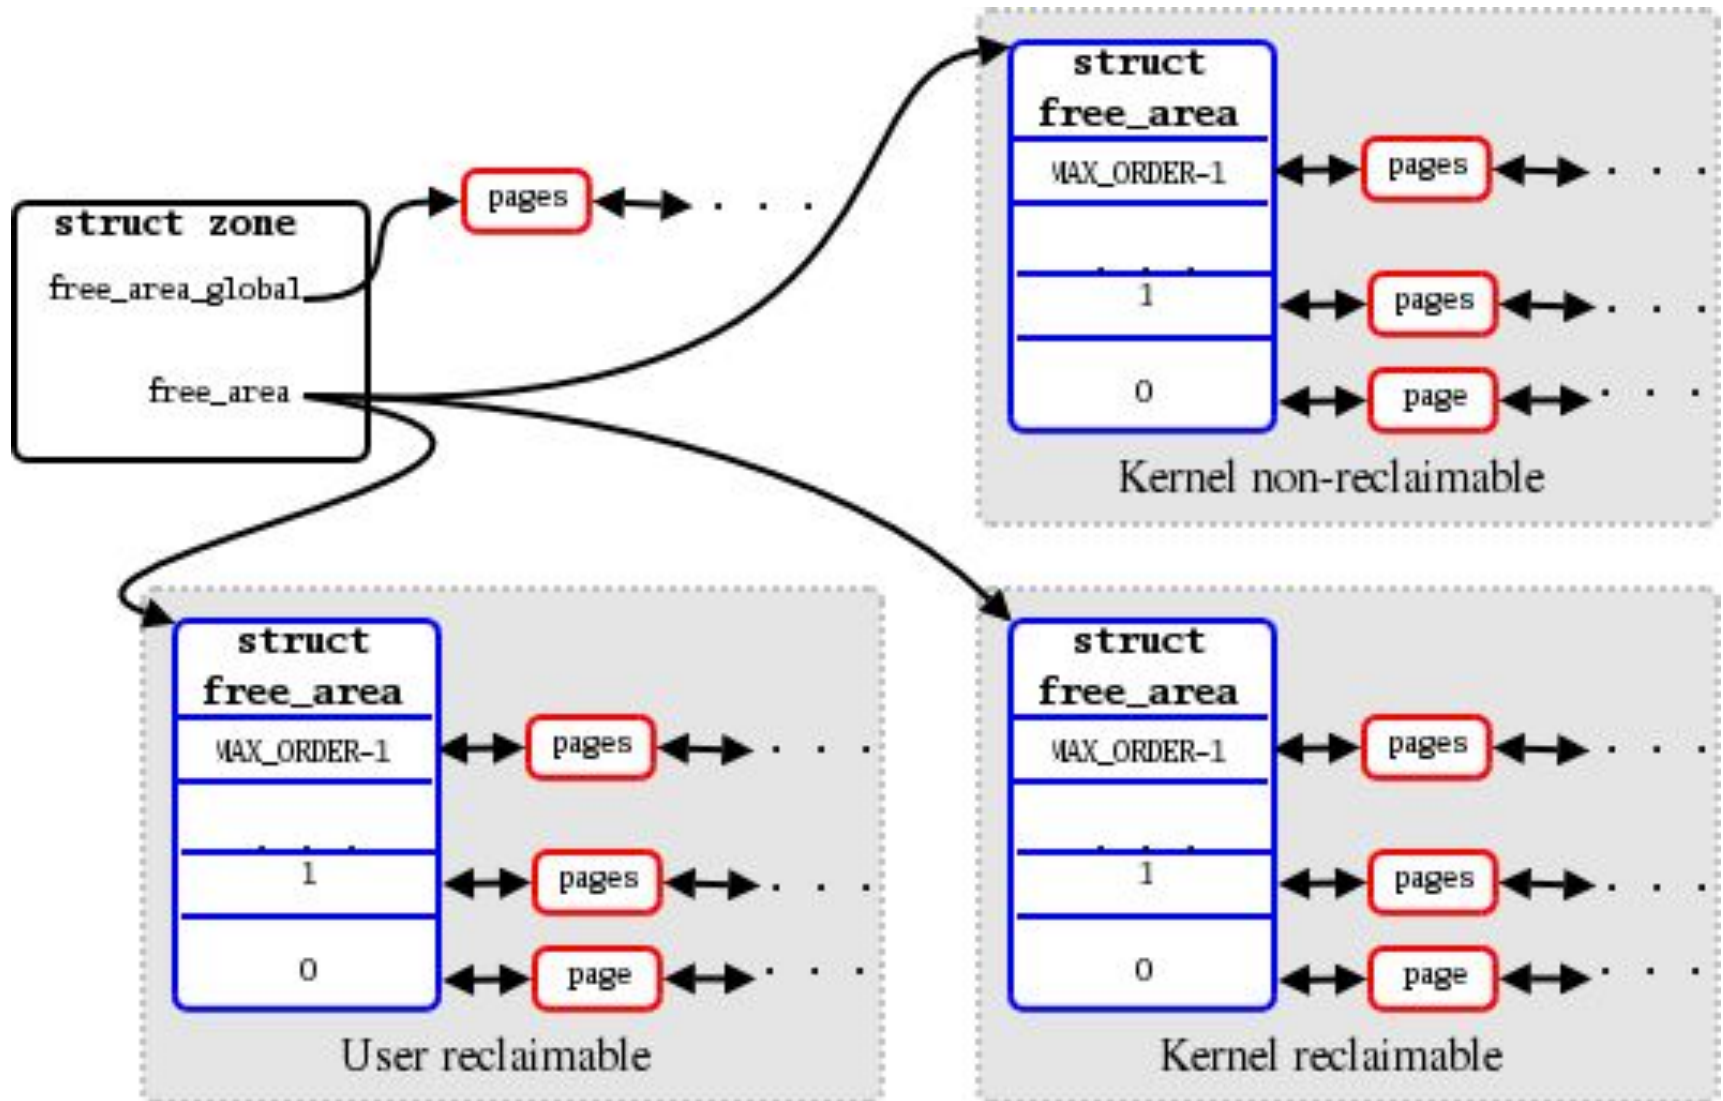
Buddy addresses



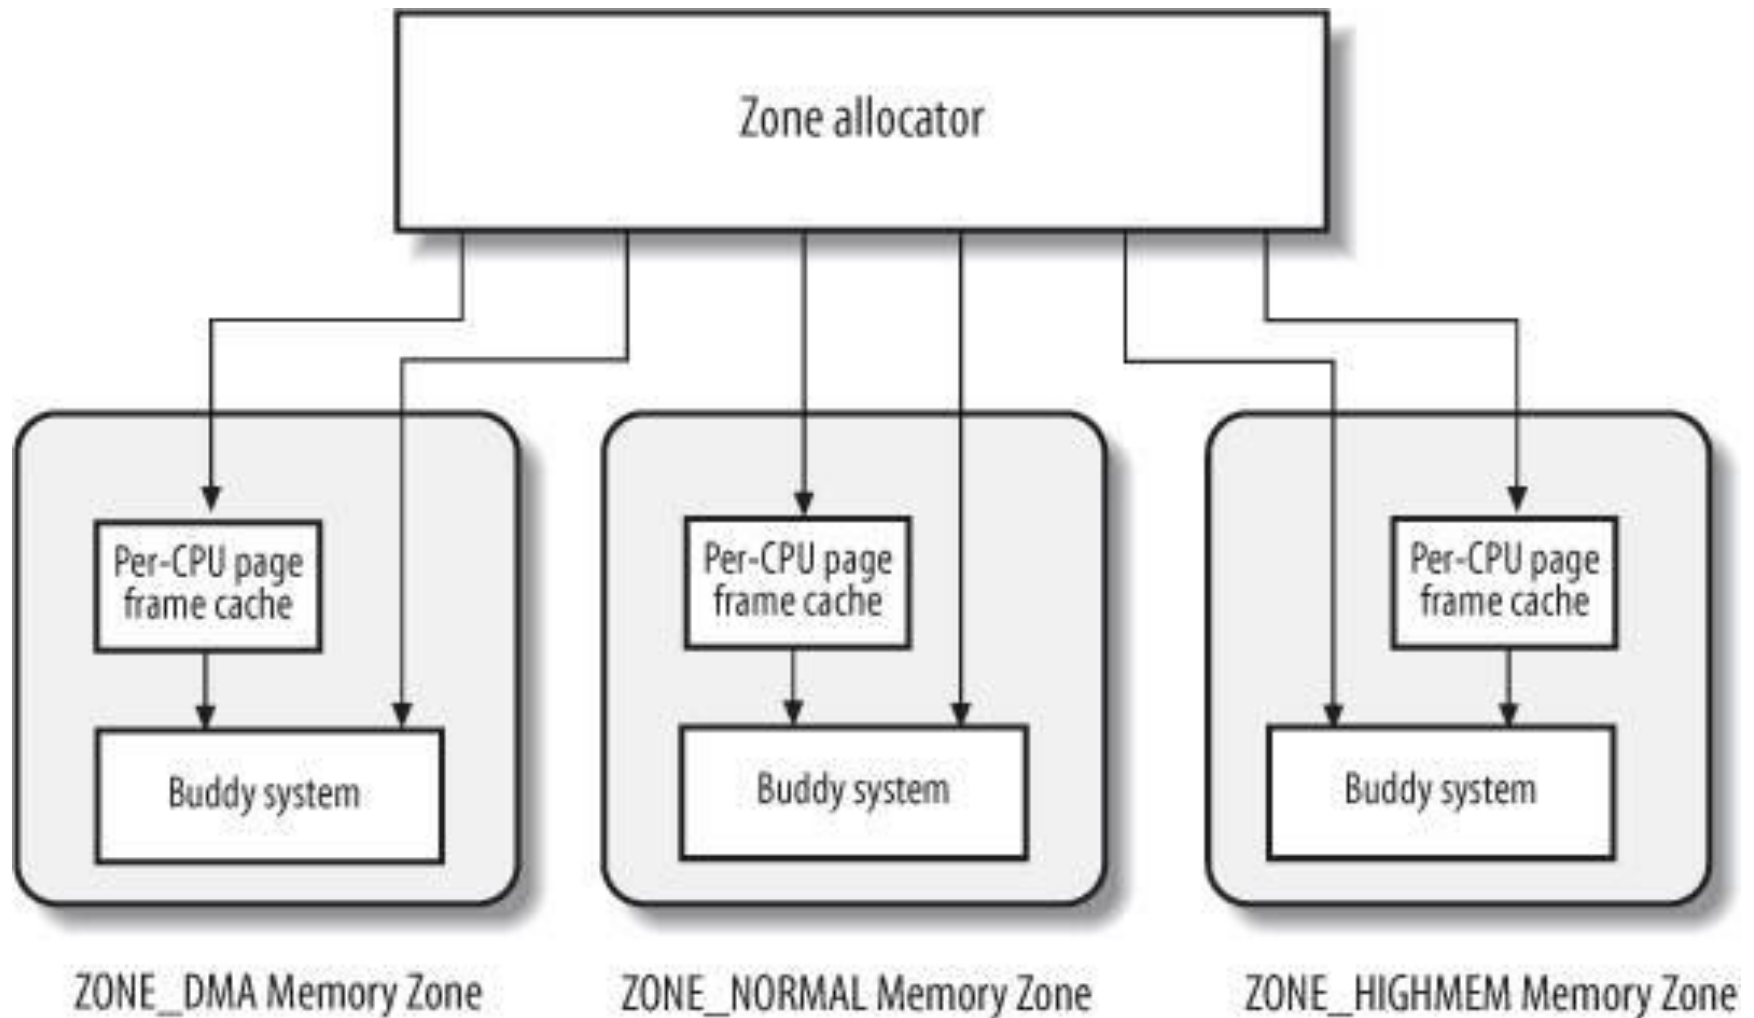
The Linux Buddy Allocator: Buddy Operations



The Linux Buddy Allocator: Per-zone Data Structures



The Linux Buddy Allocator: Per-zone Data Structures



The Linux Buddy Allocator: Per-CPU Frame Cache

The Linux Buddy Allocator: Interface

- Core API (with many wrappers):
 - `struct page * alloc_pages(gfp_t gfp_mask, unsigned int order);`
 - `void __free_pages(struct page *page, unsigned int order);`
- GFP (Get Free Page) flags (`linux/gfp.h`):

```
19  #define      ___GFP_DMA                0x01u
20  #define      ___GFP_HIGHMEM            0x02u
21  #define      ___GFP_DMA32              0x04u
22  #define      ___GFP_MOVABLE            0x08u
24  #define      ___GFP_HIGH                0x20u
30  #define      ___GFP_NOFAIL              0x800u
34  #define      ___GFP_ZERO                0x8000u
38  #define      ___GFP_ATOMIC              0x80000u
```

From Linux to Lab1 in OpenLSD

- UMA (1 node)
- Single memory type (1 zone)
- Simple alloc-only boot-time allocator
- Simple buddy page frame allocator

The OpenLSD Boot Memory Allocator

- A simple page-granular buffer allocator
 - `boot_alloc(uint32_t n);`
- Maintains an index to the next free block
- Can't free or reuse memory
- What happens when out of memory?
- Can't hand over memory to other allocators
- Only used for very few long-lived data structures allocated early on
 - e.g., initial page directory

The OpenLSD Page Frame Allocator

- A simple buddy page frame allocator
 - `struct page_info *page_alloc(int flags);`
 - `void page_free(struct page_info *pp);`
- No built-in sanity checks
- Bonus: Linux-like sanity checks (details later)
- Can allocate one page at the time
- How many freelists of page descriptors?
 - `page_free_list[]` (kernel/mem/buddy.c)


The Linux Buddy Allocator: Fragmentation

- External fragmentation:
 - Inability to service allocations due to excessive number of small free blocks
 - How addressed:
 - free coalescing
 - `vmalloc` allocator for large allocations
- Internal fragmentation:
 - Space wastage due to large block assigned to smaller allocation(s)
 - How addressed:
 - Slab allocators

The Linux `vmalloc` allocator

- Allocator for large allocations to deal with external fragmentation:
 - `void *vmalloc(unsigned long size);`
 - `void vfree(const void *addr);`
- Basic behavior:
 - First-fit, similar in spirit to `memblock`
- To satisfy allocation of size `N`:
 - Allocate `N` page frames
 - Map page frames in virtually continuous buffer

The Linux vmalloc allocator



The Linux vmalloc allocator

mm/vmalloc.c:


```
static void *__vmalloc_area_node(struct vm_struct *area, gfp_t gfp,
                                pgprot_t prot, int node)
{
    for (i = 0; i < area->nr_pages; i++) {
        struct page *page;
        page = alloc_pages_node(node, alloc_mask|highmem_mask, 0);
        area->pages[i] = page;
    }

    if (map_kernel_range(area->addr, get_vm_area_size(area),
                        prot, pages)<0)
        goto fail;
    return area->addr;
}
```

The Linux slab allocator

- Allocator for small allocations to deal with internal fragmentation (high-level interface):
 - `void *kmalloc(size_t size, gfp_t flags);`
 - `void kfree(const void *addr);`
- Unlike `vmalloc`:
 - Many implementations:
 - SLAB, SLOB, SLUB, ...
 - Memory is allocated from per-object-size *caches*:
 - Typical range [8B; 8KB]
 - Allocated memory is physically contiguous


The Linux slab allocator




The Linux slab allocator

- Cache interface:

- `struct kmem_cache *kmem_cache_create(...);`
- `void kmem_cache_destroy(struct kmem_cache*);`
- `void *kmem_cache_alloc(struct kmem_cache*,...);`
- `void kmem_cache_free(struct kmem_cache*,...);`




The Linux slab allocator

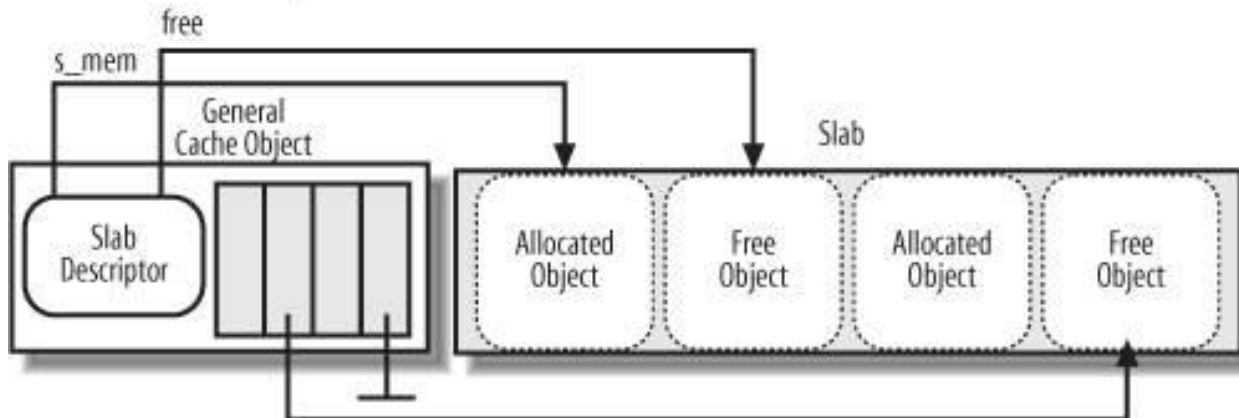



The Linux slab allocator

Slab with Internal Descriptors



Slab with External Descriptors





Frame Allocation: Memory Errors

Frame Allocation: Sanity Checks


- CONFIG_DEBUG_PAGEALLOC
- 1 - Out-of-bounds detection
 - **Page guarding** (page_is_guard)



- **Properties:**
 - Immediately detects each out-of-bounds access
 - Only accesses to [-4KB, +4KB]

Frame Allocation: Sanity Checks


- CONFIG_DEBUG_PAGEALLOC
- 1 - Out-of-bounds detection
 - **Page canaries** (unimplemented alternative)



- **Properties:**
 - Lazily detects out-of-bounds writes by mismatch
 - Only writes to $[-\text{sizeof}(\text{canary}), +\text{sizeof}(\text{canary})]$

Frame Allocation: Sanity Checks


- CONFIG_DEBUG_PAGEALLOC
- 2 - Use-after-free detection
 - **Page remapping** (`kernel_map_pages`)



- **Properties:**
 - Detects each use-after-free when unallocated
 - What happens after memory reuse?

Frame Allocation: Sanity Checks

- CONFIG_DEBUG_PAGEALLOC
- 2 - Use-after-free detection
 - **Page poisoning** (`kernel_poison_pages`)



- **Properties:**
 - Lazily detects each write-after-free by mismatch
 - Can't handle memory reuse

Frame Allocation: Sanity Checks

- A number of object-level *sanitizers*
- `kmemcheck` (uninitialized reads, now off-tree)
 - Memory protection-based
 - Traps at every read to check if data initialized
- `kmemleak` (memory leaks)
 - Periodic conservative garbage collection
 - Reports objects not pointed by any likely pointers
- `kasan` (out-of-bounds and use-after-free)
 - Compiler-based
 - Combines canary and poisoning ideas
 - Checks each access to check if target is live object
- Others: UBSAN, KCSAN

Frame Allocation: Sanity Checks

- Built-in integrity checks in frame allocator
- `check_new_page`
 - Semantic checks on page descriptor
 - Memory error checks?
- `free_pages (virt_addr_valid)`
 - Invalid free detection
 - Limitations?

References


- [1] Gorman, Mel. *Understanding the Linux Virtual Memory Manager*, 2007.
- [2] “Development tools for the kernel,” n.d. <https://www.kernel.org/doc/html/latest/dev-tools>.
- [3] “The ‘too Small to Fail’ memory-Allocation Rule,” n.d. <https://lwn.net/Articles/627419>.

Managing Physical Memory

Backup Slides

What is “Physical Memory” Anyway?

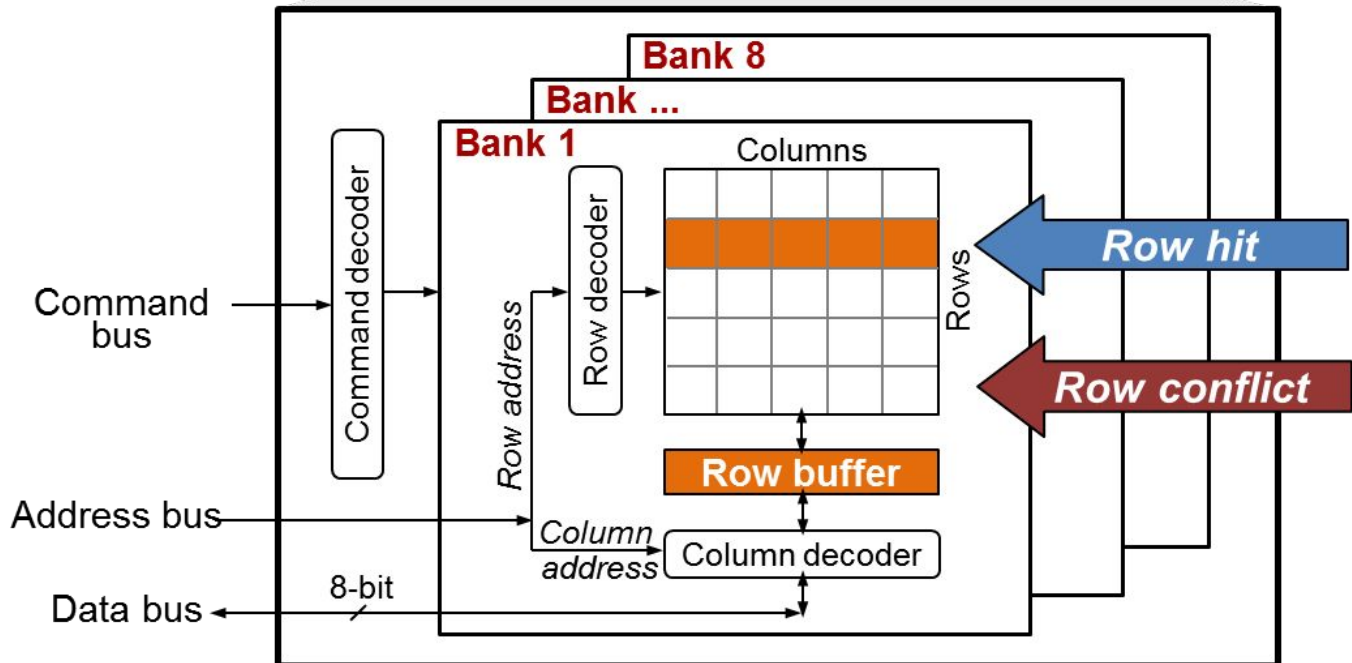
DRAM Rank



DRAM Chip

BANK example
(R x c x bits):

4096 x 1024 x 64



What is “Physical Memory” Anyway?

```
mov [0x7f00], 1 # real mode
```

Q: What gets physically written to memory?

A: Hard to say in practice:

- Address scrambling
 - Complex *paddr* to DRAM <C, R, B> mapping
 - Distributes load to avoid “bank trashing”
- Data scrambling
 - Logical 1 != DRAM 1
 - Avoids data burst (resonance) on the data bus
 - Mitigates (some) cold-boot attacks

What is Wrong With DRAM?

Side channels:

- Bank collisions can be detected via timing
- Allows attackers to detect sensitive events

Hardware faults:

- Capacitors can drop charge too quickly
- Bit flips due to cross-cell/-row interference
 - **Rowhammer**: originally reliability problem
 - **Flip Feng Shui**: bit flips controllable and exploitable
 - ECC *might* make this better

What is Wrong With DRAM?

**You will learn more about this
at “Hardware Security”...**