

Multiprocessing

Advanced Operating Systems

Overview

- Process creation
 - fork
 - exec
- Scheduling
 - Time management
 - Scheduling strategies
- IPC
 - Shared memory
 - Semaphores
 - Message queues

Fork

```
pid_t fork(void);
```

- Creates a new child process by duplicating the calling process
- Implemented on top of the `clone()` system call on Linux

```
kernel/fork.c:
```

```
/*  
 * Fork is rather simple, but the memory  
 * management can be a bitch. See 'copy_page_range()'   
 */
```

Fork: The Simple Part

- Duplicate task (`dup_task_struct`):
 - Copy most information, with exceptions, e.g. PID
- Allocate and initialize new kernel stack:
 - Setup `thread_info`
 - Copy trap frame (`pt_regs`) and update (e.g., `%rax`)
- And more:

```
retval = copy_files(clone_flags, p);  
retval = copy_fs(clone_flags, p);  
retval = copy_sighand(clone_flags, p);  
retval = copy_signal(clone_flags, p);  
retval = copy_mm(clone_flags, p);
```

Fork: The copy_mm Part

- Duplicate mm descriptor (dup_mm):
 - Copy over basic information
 - Initialize empty address space (new page directory)
- Duplicate address space (dup_mmap):
 - Copy over VMA information
 - Copy page tables (copy_page_range)
 - Fixup page table entries
 - MAP_SHARED VMAs: Share page frames
 - MAP_PRIVATE (R) VMAs: Share page frames
 - MAP_PRIVATE (R/W) VMAs: COW page frames


Fork: The copy_mm Part

```
/*
 * If it's a COW mapping, write protect it both
 * in the parent and the child
 */
if (is_cow_mapping(vm_flags)) {
    ptep_set_wrprotect(src_mm, addr, src_pte);
    pte = pte_wrprotect(pte);
}

/*
 * If it's a shared mapping, mark it clean in
 * the child
 */
if (vm_flags & VM_SHARED)
    pte = pte_mkclean(pte);

page = vm_normal_page(vma, addr, pte);
get_page(page);
set_pte_at(dst_mm, addr, dst_pte, pte);
```

Fork: The copy_mm Part



- Each common page frame: 2 VMAs, 2 PTEs
- Shmem: $\text{perm}(\text{VMA}) = \text{perm}(\text{PTE})$
- COW: $\text{perm}(\text{VMA}) = \text{R/W}$, $\text{perm}(\text{PTE}) = \text{R}$
 - Page fault handler recognizes COW (`do_wp_page`)
 - Duplicates page frame and remaps a new private copy into the faulting PTE on demand

Exec

```
int exec*(const char *filename, ...);
```

- Executes the program pointed by filename
- Implementation:
 - Input and permission checking
 - Load binary headers in memory (prepare_binprm)
 - Find the binary format (search_binary_handler)
 - Flush old resources:
 - Reinitialize task_struct and (empty) mm
 - Flush VMAs, page tables, page frames
 - Load the binary (load_binary)

Exec

- Load binary (binary-format specific):
 - Parse headers and sections
 - Create corresponding VMAs (Data, Text, Stack, etc.)
 - Update %rsp and %rip in trap frame
 - %rsp = top of the user stack
 - %rip = program entry point for statically linked binaries
 - %rip = dynamic linker's entry point otherwise
- Page tables initially empty
 - Even binary file pages (e.g., text) demand paged
 - Page fault handler maps them from page cache using COW-based strategy for safe sharing

Copy-On-Write: Applications

- MAP_PRIVATE file pages
 - Deduplicates binary pages for unrelated processes
 - Many (e.g., text pages) never COWed
- MAP_PRIVATE anon forked pages
 - Deduplicates pages within process hierarchy
 - Many (e.g., fork+exec) never COWed
- MAP_PRIVATE anon zero pages
 - Deduplicates zero pages for unrelated processes
 - At first read PF, map single read-only zero page
 - COW at first write PF

Scheduling

- We can now create multiple processes, how do we schedule them on a given CPU?
- The easy version (OpenLSD):
 - Tell the hardware to raise periodic timer interrupts
 - (Timer) interrupt handler invokes simple scheduler
- Simple scheduler:
 - Maintains a scheduling queue (FIFO)
 - Enqueues interrupted process at tail
 - Dequeues process at head and runs it on CPU
- This is a very simple preemptive round-robin scheduler, with no priorities, no fairness, etc.

Multiprocessing in OpenLSD

Lab 5: All together now (multiprocessing)

Core:

- Support multiprocessing with simple round-robin scheduling
- Support fork with Copy-On-Write (COW)

Bonuses:

- exec, more COW, IPC, timers

Time Management

- Hardware offers clock and timer circuits
- Clock circuits:
 - Expose counters incremented at given frequency
 - Can be used to keep track of current time of day
 - Can be used for precise time measurements
- Timer circuits:
 - Issue periodic interrupts at given frequency
 - Can be used for scheduling
 - Can be used to implement kernel and user timers

Time Management: Linux Clock Sources

- Abstraction of a clock circuit (clocksource)
\$ cat
/sys/devices/system/clocksource/clocksource0/
available_clocksource
tsc hpet acpi_pm
- Clock sources (with different precision):
 - tsc: Timestamp counter (default)
 - hpet: high-precision event timer
 - acpi_pm: ACPI power management timer

Time Management: Real Time Clock (RTC)

- Special “source” (not a `clocksource`)
 - Persistent (battery-powered)
 - Low-precision (seconds)
- Used to get the current date and time:

```
$ cat /sys/class/rtc/rtc0/date
```

```
2016-09-26
```

```
$ cat /sys/class/rtc/rtc0/time
```

```
15:03:00
```

Time Management: Linux Clock Event Devices

- Abstraction of a timer circuit
(`clock_event_device`)

```
$ cat /proc/timer_list | grep ^Clock | sort -u
```

```
Clock Event Device: hpet
```

```
Clock Event Device: lapic #local APIC
```
- Clock event device programmed to issue interrupts at `CONFIG_HZ` frequency, e.g.:
 - `CONFIG_HZ=250`, timer interrupt (or *tick*) every 4ms
- Both user and kernel preemption are possible
 - OpenLSD: non-preemptive kernel

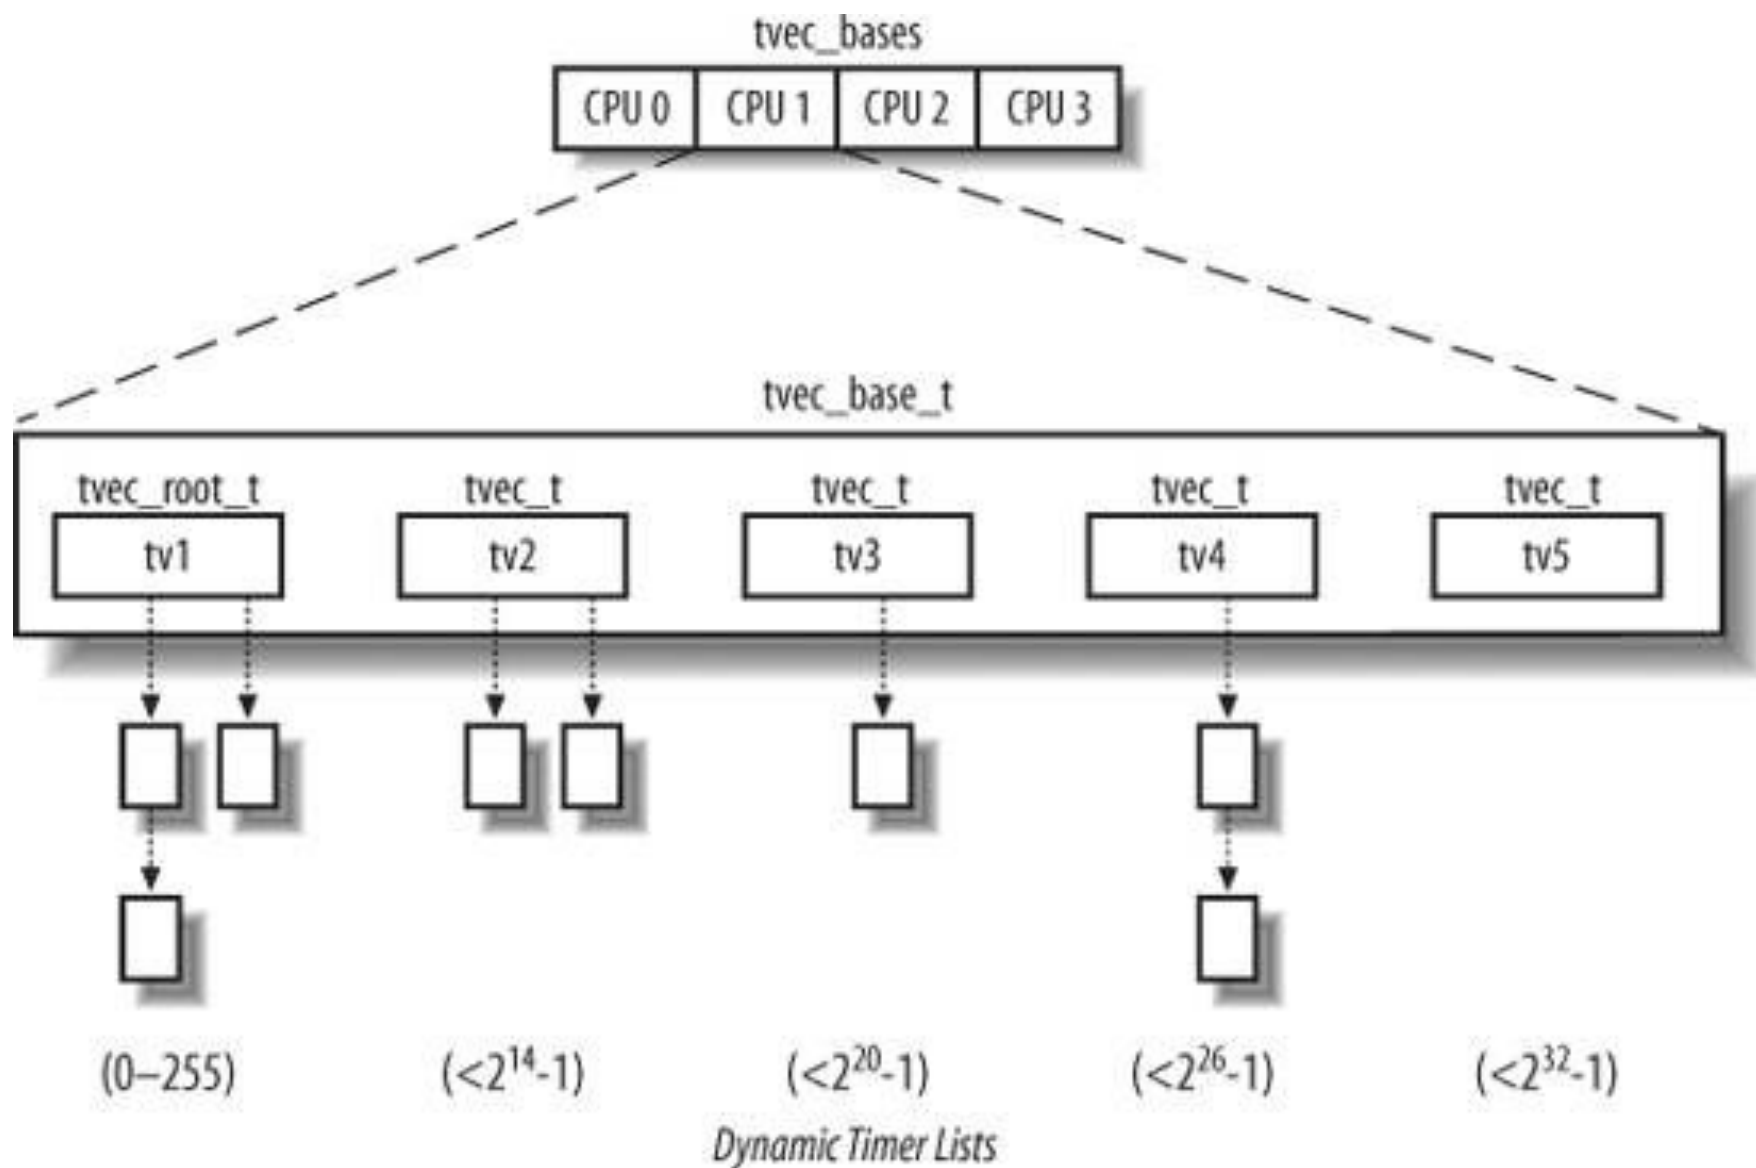
Time Management: Linux Timer Interrupts

- At each tick (roughly):

```
jiffies_64++; // update ticks since startup
update_wall_time(); // update current date/time
update_process_times(); // accounting
profile_tick(); // profiling
```

```
while (time_after_eq(jiffies, base->clk))
    expire_timers(base, ...); // expired timers
```

```
schedule(); // invoke the scheduler
```



Linux Software Timers

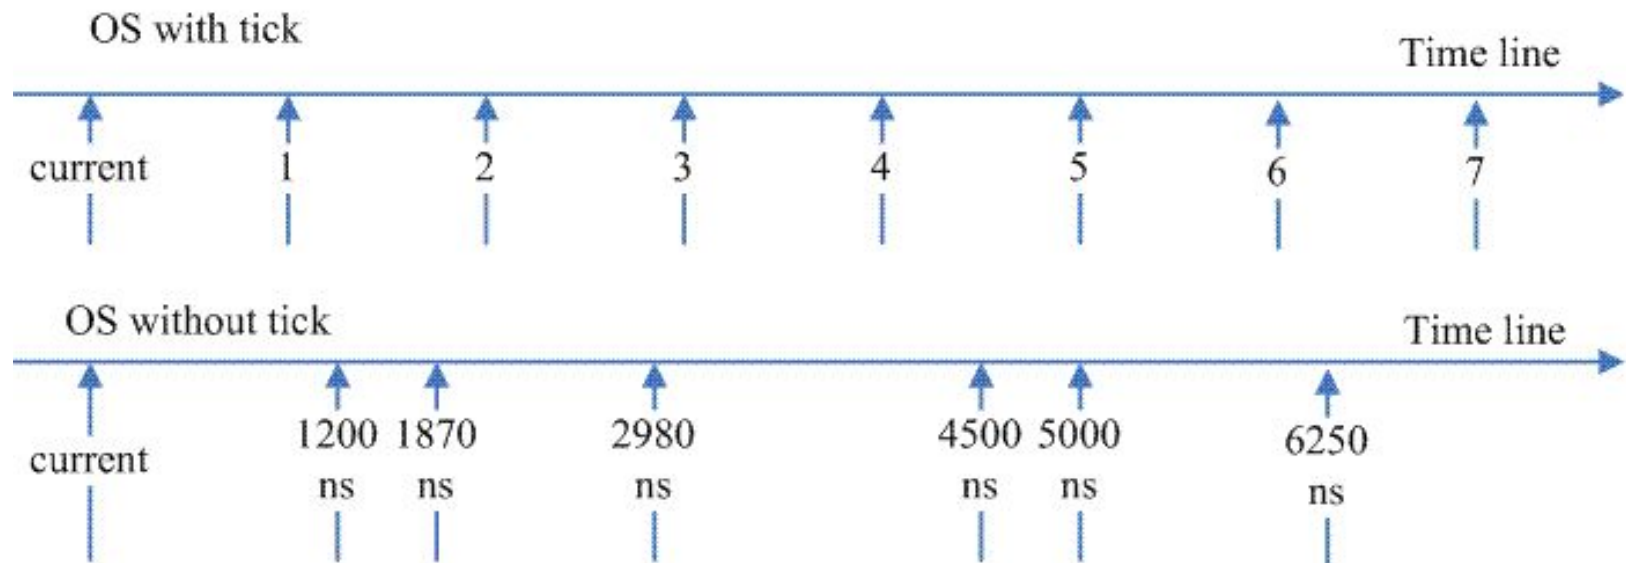
Scheduling:

Per-task Building Blocks

- State: Running, Runnable, Sleeping
- Quantum or time slice
 - Max number of jiffies a task can run on CPU
 - Initialized with task-specific formula (e.g., 100ms)
 - Decrement at every tick, task is done at 0
 - Sufficient to ensure fairness for CPU-intensive tasks
- Priority
 - Initialized with static (predetermined) priority
 - Possibly adjusted periodically
- Scheduling policy
 - NORMAL, BATCH, IDLE → Completely fair scheduler
 - FIFO, RR → Real-time scheduler

Scheduling: Tickless Kernel

- Can we go tickless (+responsive, -overhead)?
 - Create a new sw timer for “end of quantum” event
 - Reprogram hardware to tick at next timer to expire
- NO_HZ_IDLE|FULL: Tickless when idle|1 task



Scheduling:

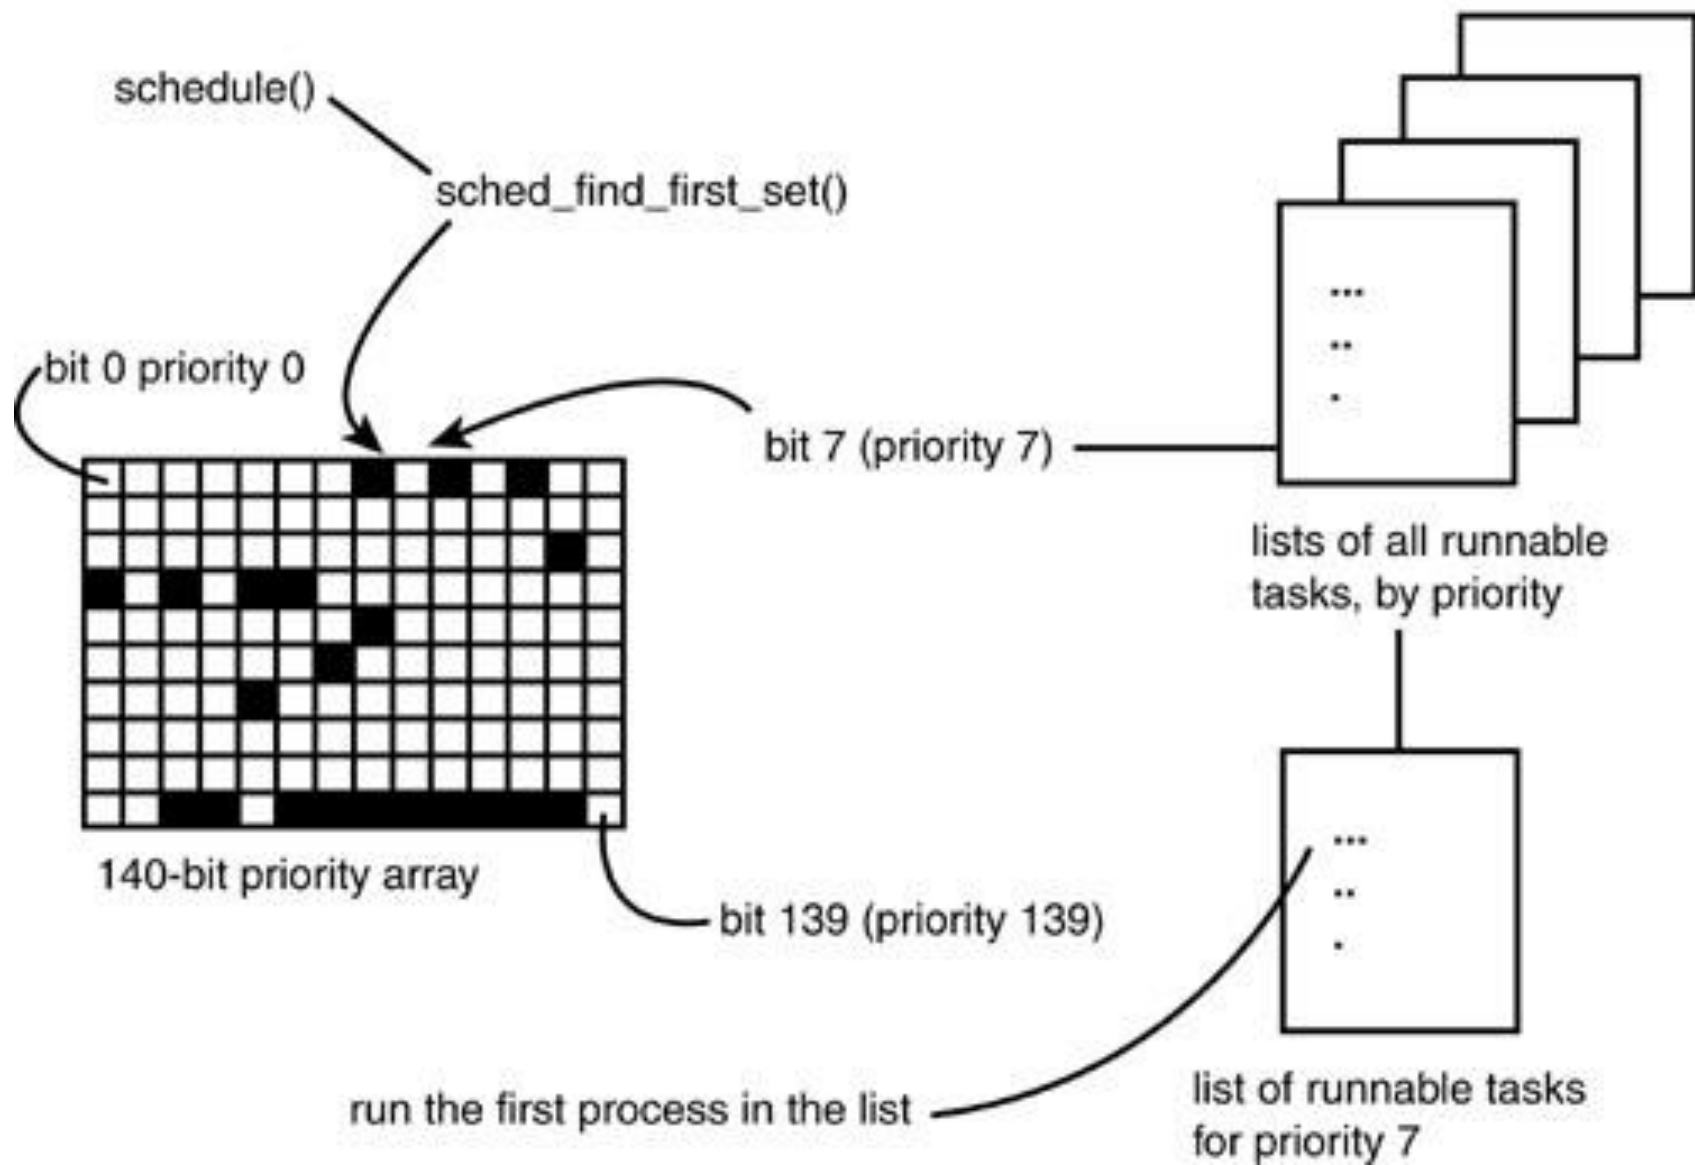
Linux O(1) Scheduler

- Preemptive round-robin priority scheduler
- Once default, now close to RT scheduler
- Maintains N run queues (1 per priority level)
- Scheduling strategy:
 - Find the highest-priority queue with runnable task
 - Find the first task on that queue and dequeue it
 - Calculate its time slice size based on priority and run it
 - When its time's up, enqueue it, and repeat
- Improving fairness:
 - Priorities are adjusted based on *sleep time*
 - Bonuses for I/O- vs. CPU-bound processes

Scheduling:

Linux O(1) Scheduler

- Why O(1) Scheduler?
 - All the operations are O(1)
 - No loop across tasks like older O(n) scheduler
- Tasks only touched at dequeue/enqueue time
 - Recalculate priorities, quanta, bonuses
- Finding the right run queue is also O(1):
 - Bitmap indicating which queues have runnable tasks
 - Find the first bit that is set (first 1-bit instruction)
 - Time depends on the number of priority levels, not on the number of tasks (*scalability*)

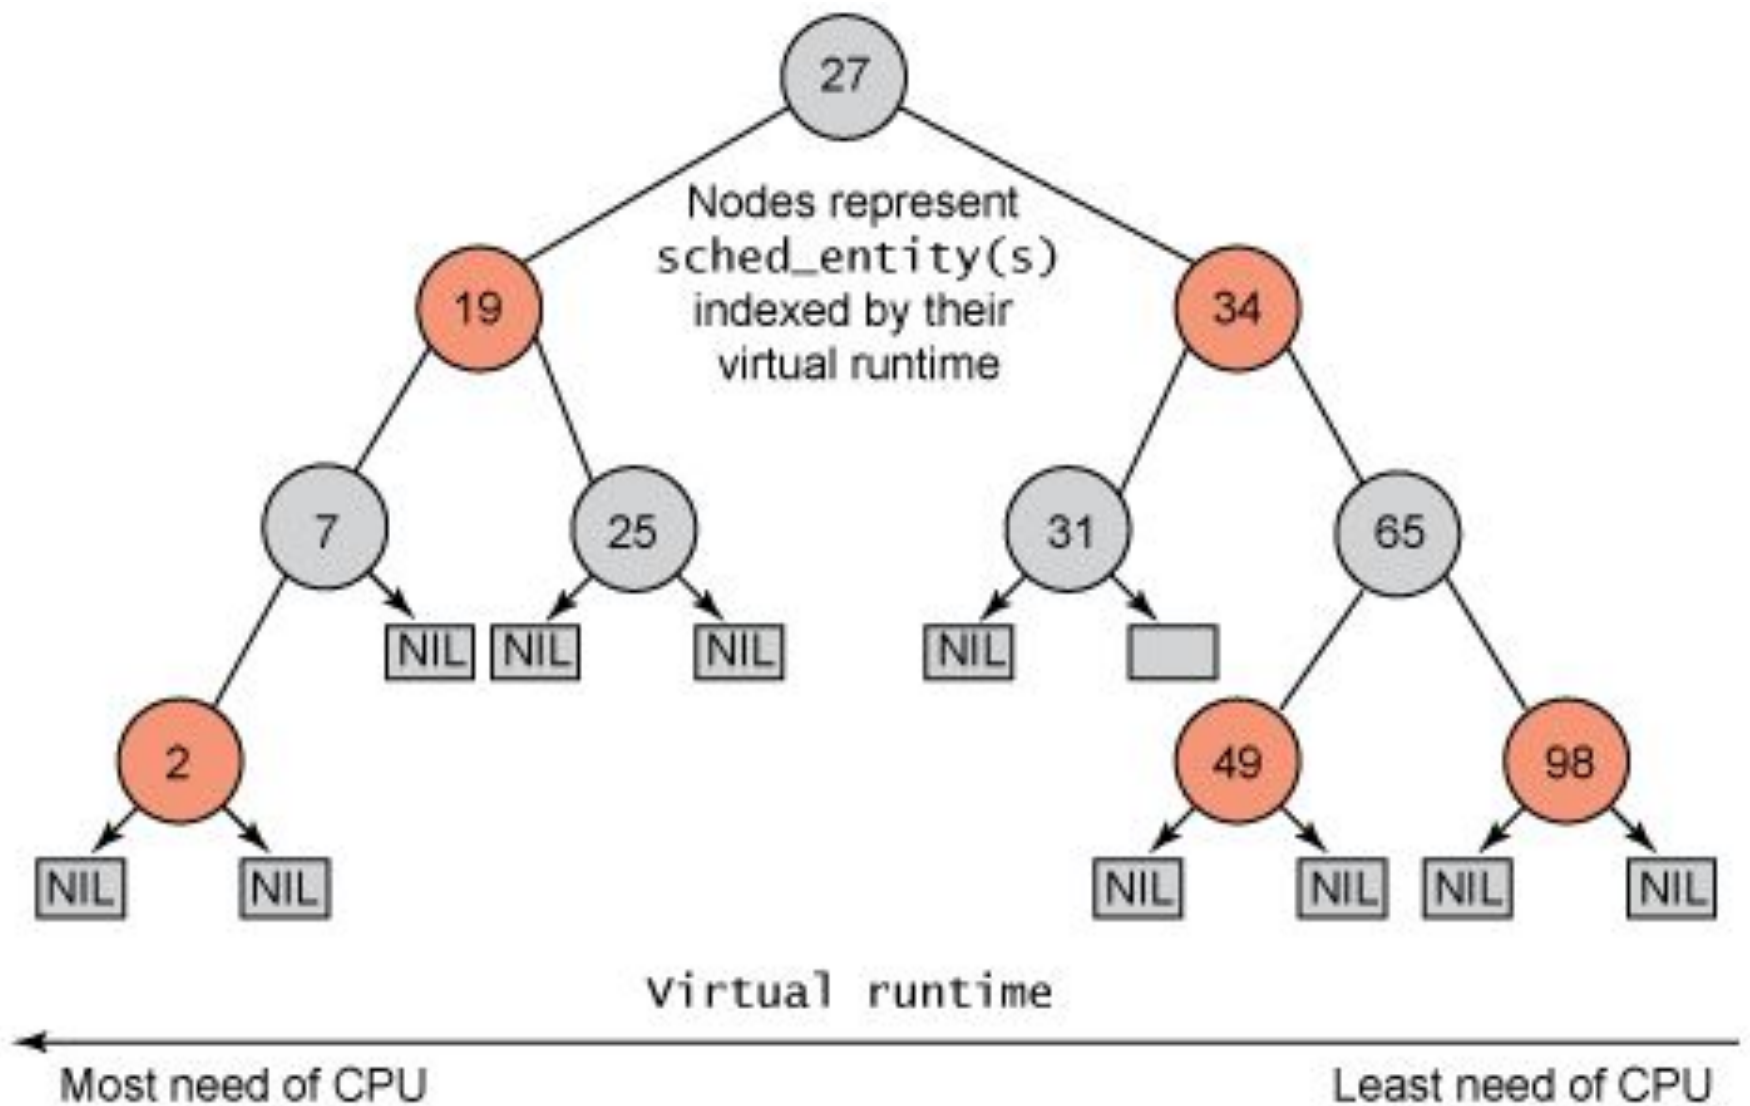


Linux O(1) Scheduler

Scheduling:

Linux CFS Scheduler

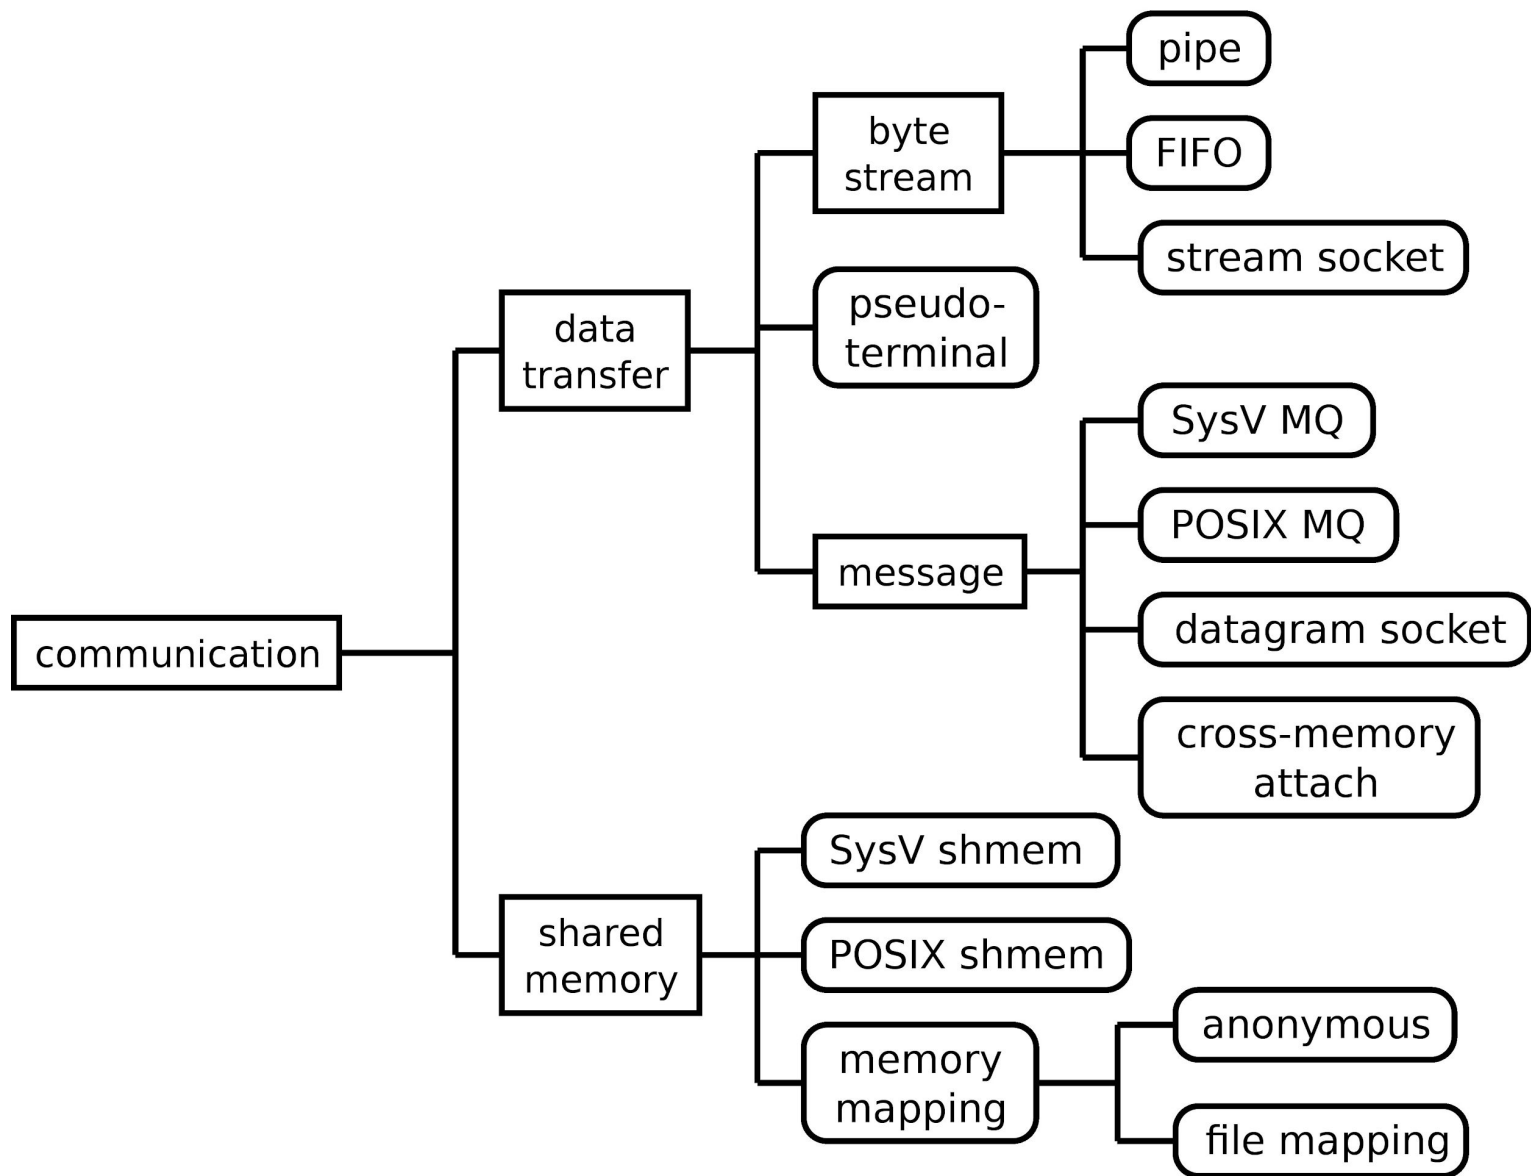
- $O(1)$ used hard-to-maintain hacks for fairness
- CFS: Tasks get a “completely fair” CPU share
- Models an “ideal, precise, multi-tasking” CPU
- Basic idea with N tasks:
 - Record how much CPU time each task has been given
 - Schedule task with biggest delta to $\text{tot_CPU_time}/N$
 - Virtualize time (*virtual runtime*) to deal with priorities
 - Increase virtual runtime faster for lower-priority tasks
- No heuristics to distinguish tasks
- No run queues, uses a red-black tree



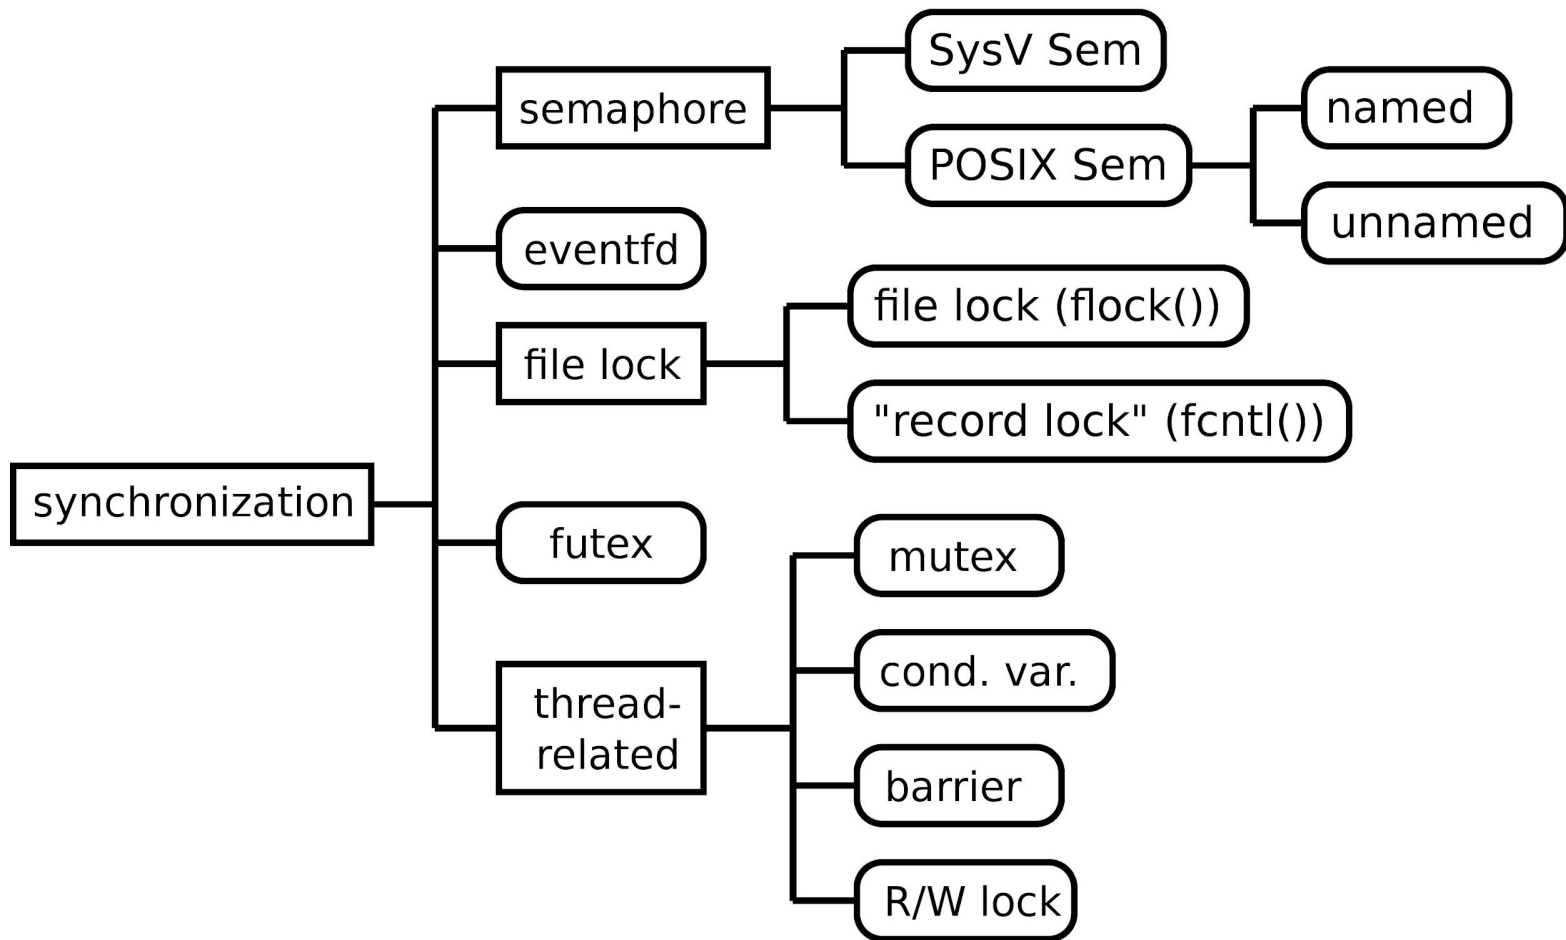
Linux CFS Scheduler

IPC

- Communication
 - Allow processes to exchange data
- Synchronization
 - Allow processes to synchronize execution
- Signals
 - Asynchronous notifications
- Multiple standards and interfaces...



IPC: Communication




IPC: Synchronization

IPC: System V vs. POSIX

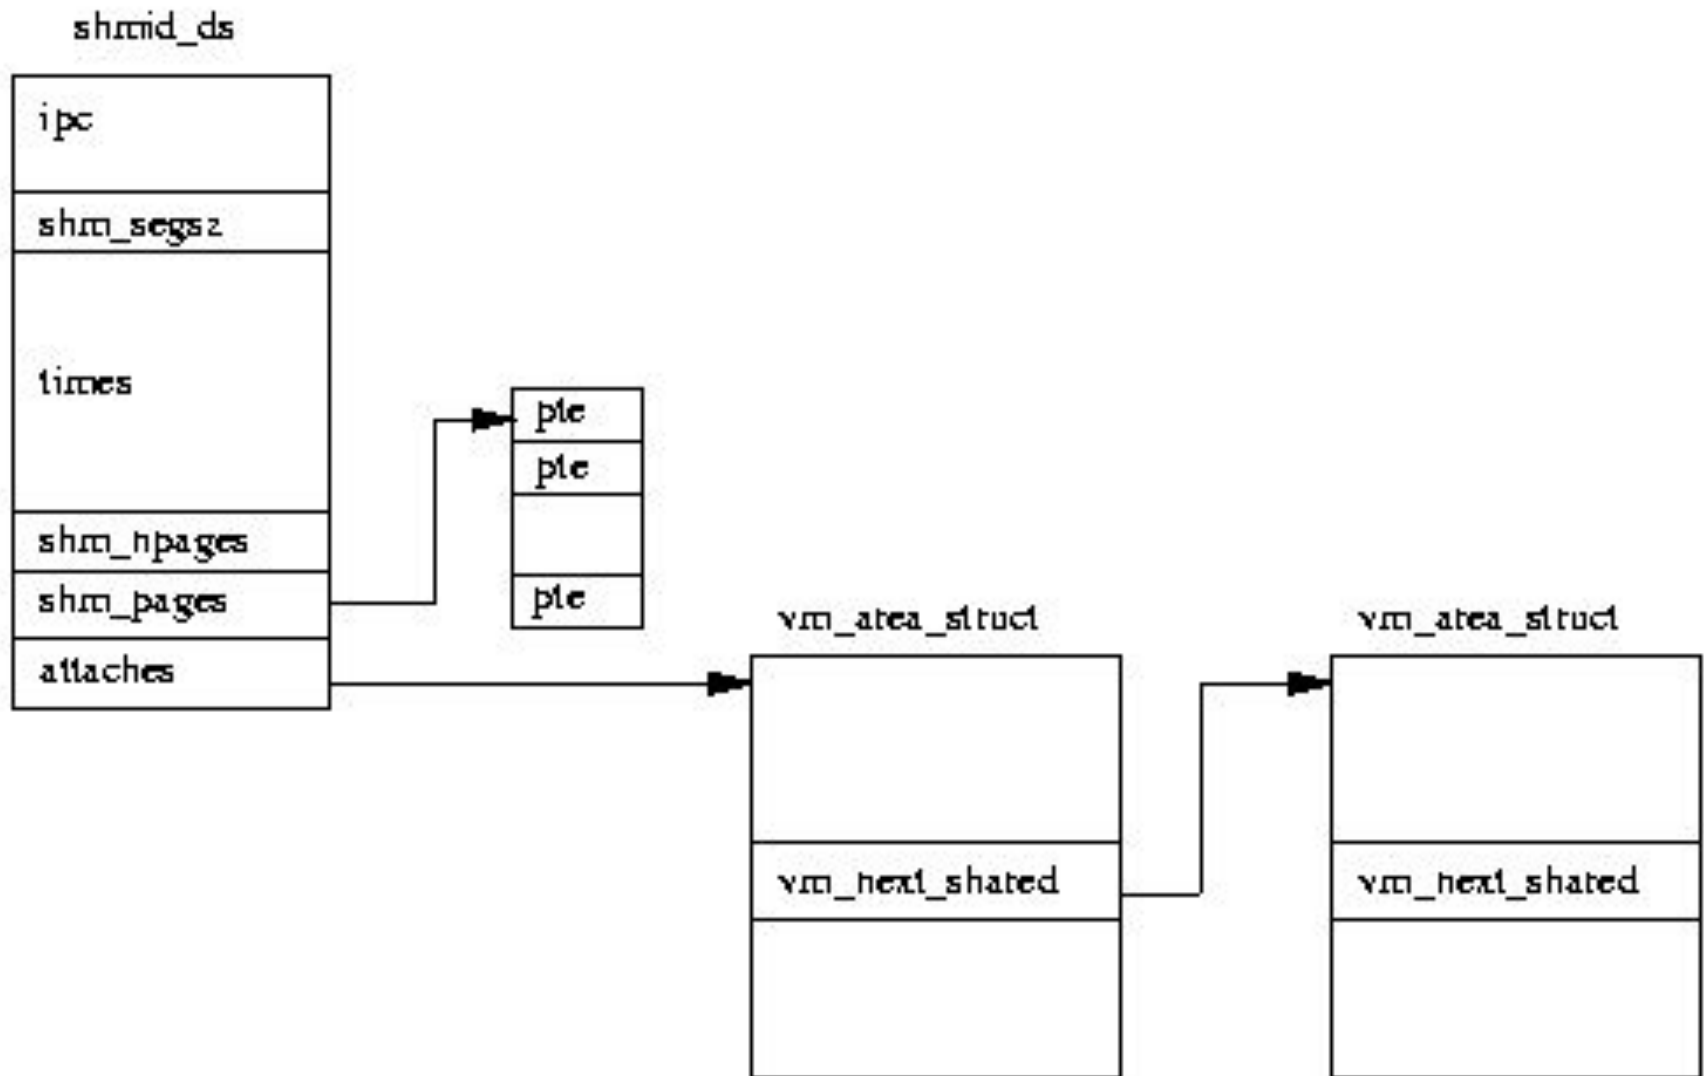
- Similar key IPC mechanisms
 - Shared memory
 - Share a region of memory
 - Semaphores
 - Synchronize with post/wait primitives
 - Message queues
 - Exchange messages between processes
- System V
 - Original UNIX implementation, +compatibility
- POSIX
 - Standardized later, +user-friendly, +features

System V: Shared Memory




System V: Shared Memory

- Get/create shmem segment by key, perm:
`int shmget(key_t key, size_t size, int shmflg);`
- Attach / detach segment shmid:
`void *shmat(int shmid, void *addr, int shmflg);`
`int shmdt(void *addr);`
- Any addr (if any) can be specified for shmat in different processes
- All the processes attaching the same segment share the same page frames



System V: Shared Memory Implementation

System V: Semaphores



System V: Semaphores

- Get/create semaphore set by key, perm:

```
int semget(key_t key, int nsems, int semflg);
```

- Perform operations on semaphore set semid:

```
int semop(int semid, struct sembuf *sops, int n);
```

```
struct sembuf {
```

```
    unsigned short sem_num;    /* number */
```

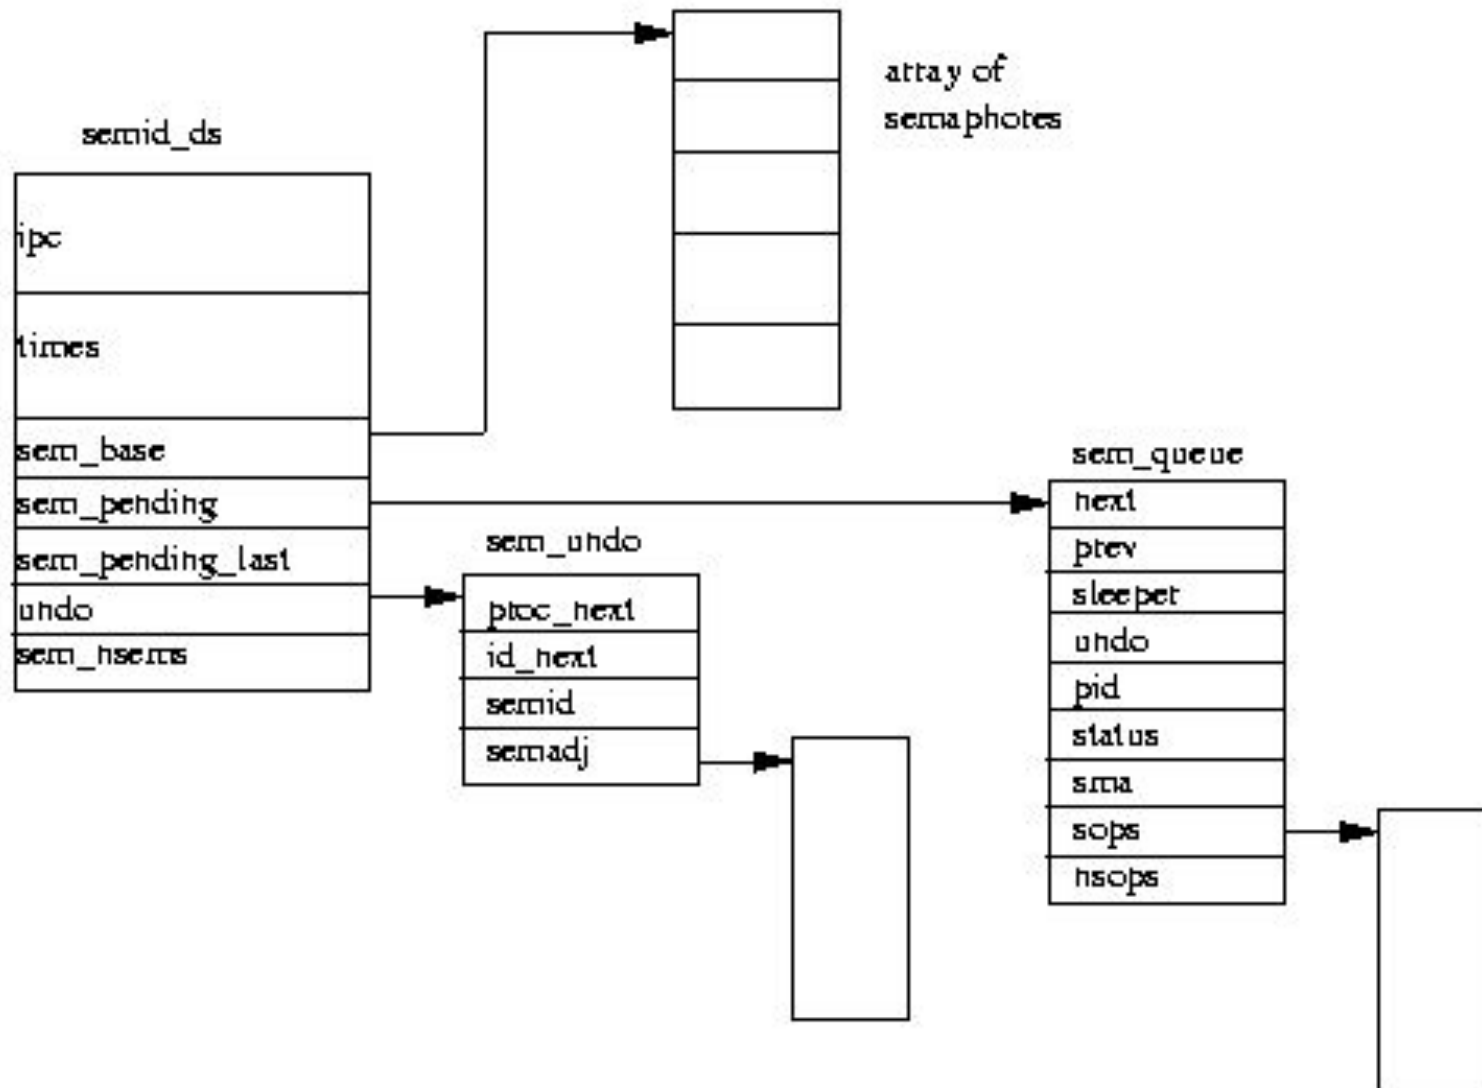
```
    short          sem_op;     /* operation */
```

```
    short          sem_flg;    /* flags */
```

```
}
```


System V: Semaphores

- Initially semaphore value (`val`) is 0
- Semaphore post (`sem_op > 0`):
 - Atomically increment semaphore by `sem_op`
- Semaphore wait (`sem_op < 0`):
 - If `val >= -sem_op`:
 - Atomically decrement semaphore by `-sem_op`
 - Else:
 - Block and wait for `val >= -sem_op`
- Semaphore waitzero (`sem_op == 0`):
 - Wait for `val == 0`



System V: Semaphores Implementation

System V: Message Queues



System V: Message Queues

- Get/create message queue by key, perm:

```
int msgget(key_t key, int msgflg);
```

- Send/receive message on message queue
msqid with an optional type:

```
int msgsnd(int msqid, const void *msgp, size_t  
msgsz, int msgflg);
```

```
ssize_t msgrcv(int msqid, void *msgp, size_t  
msgsz, long msgtyp, int msgflg);
```

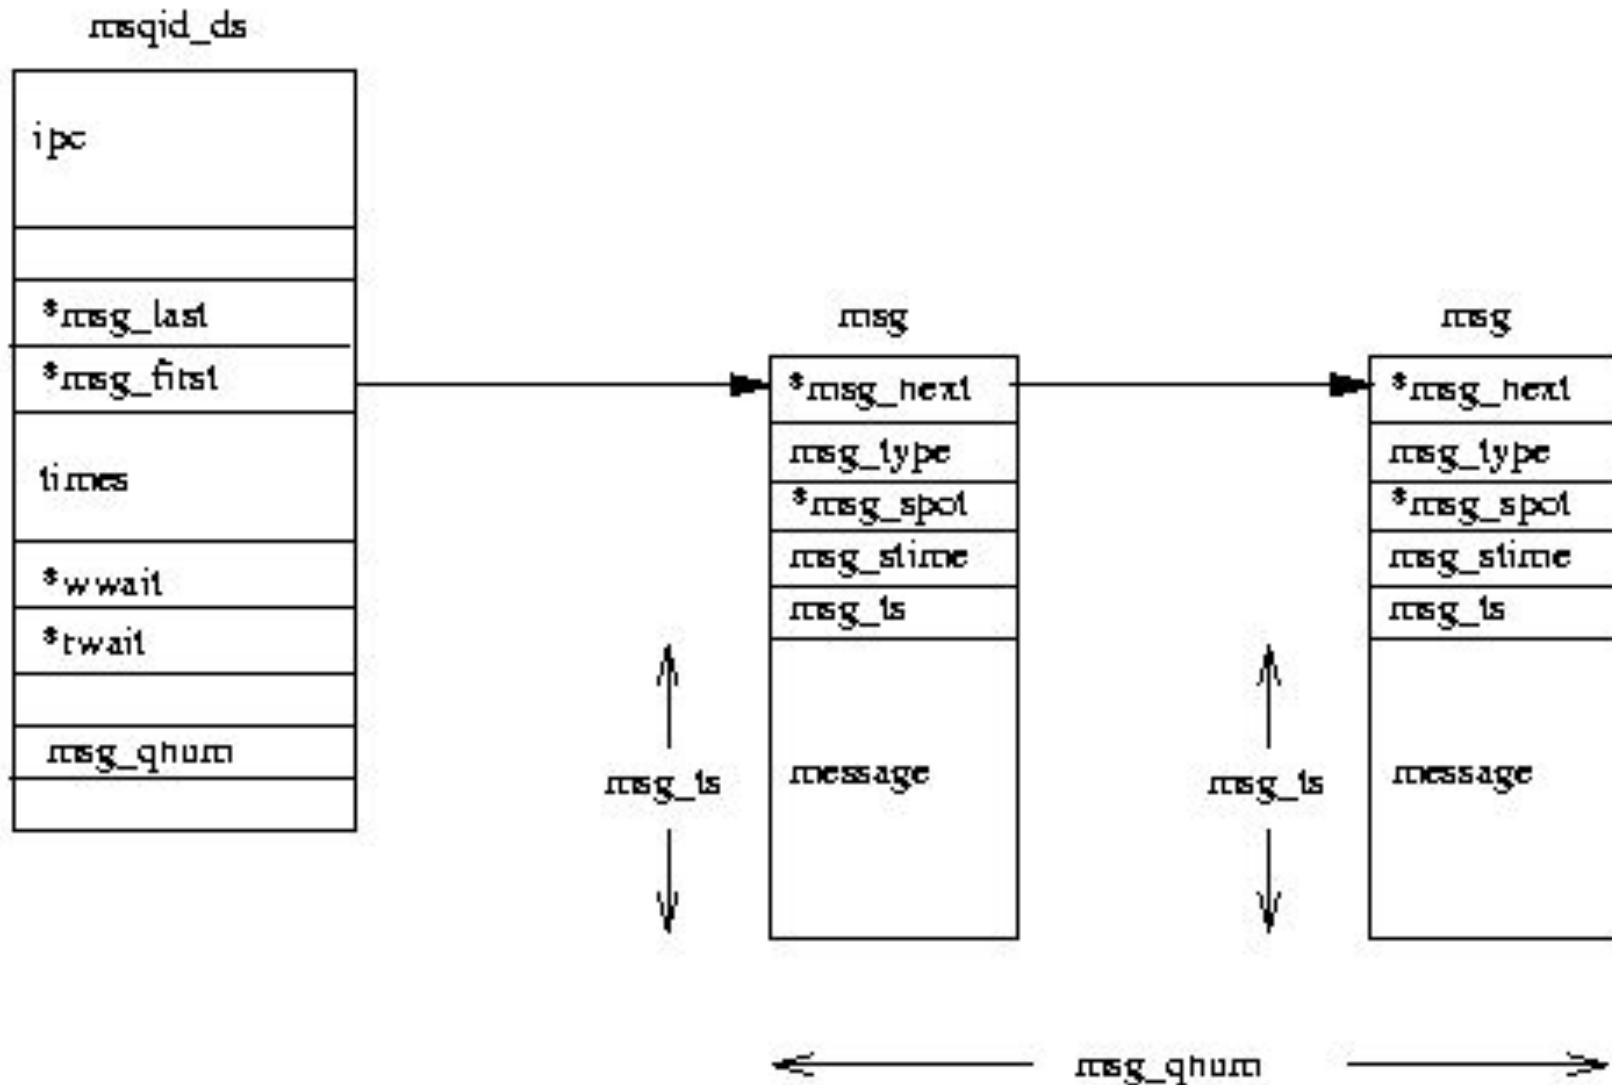
- Blocks if queue is full (send) or empty
(receive). Queue is size-limited

System V: Message Queues

- Variable-size message of fixed format:

```
struct msgbuf {  
    long mtype;    /* message type, must be > 0 */  
    char mtext[1]; /* message data */  
};
```

- Unlike pipes:
 - Not stream-oriented
 - Supports random queue access
 - Bidirectional
 - Always named



System V: Message Queue Implementation

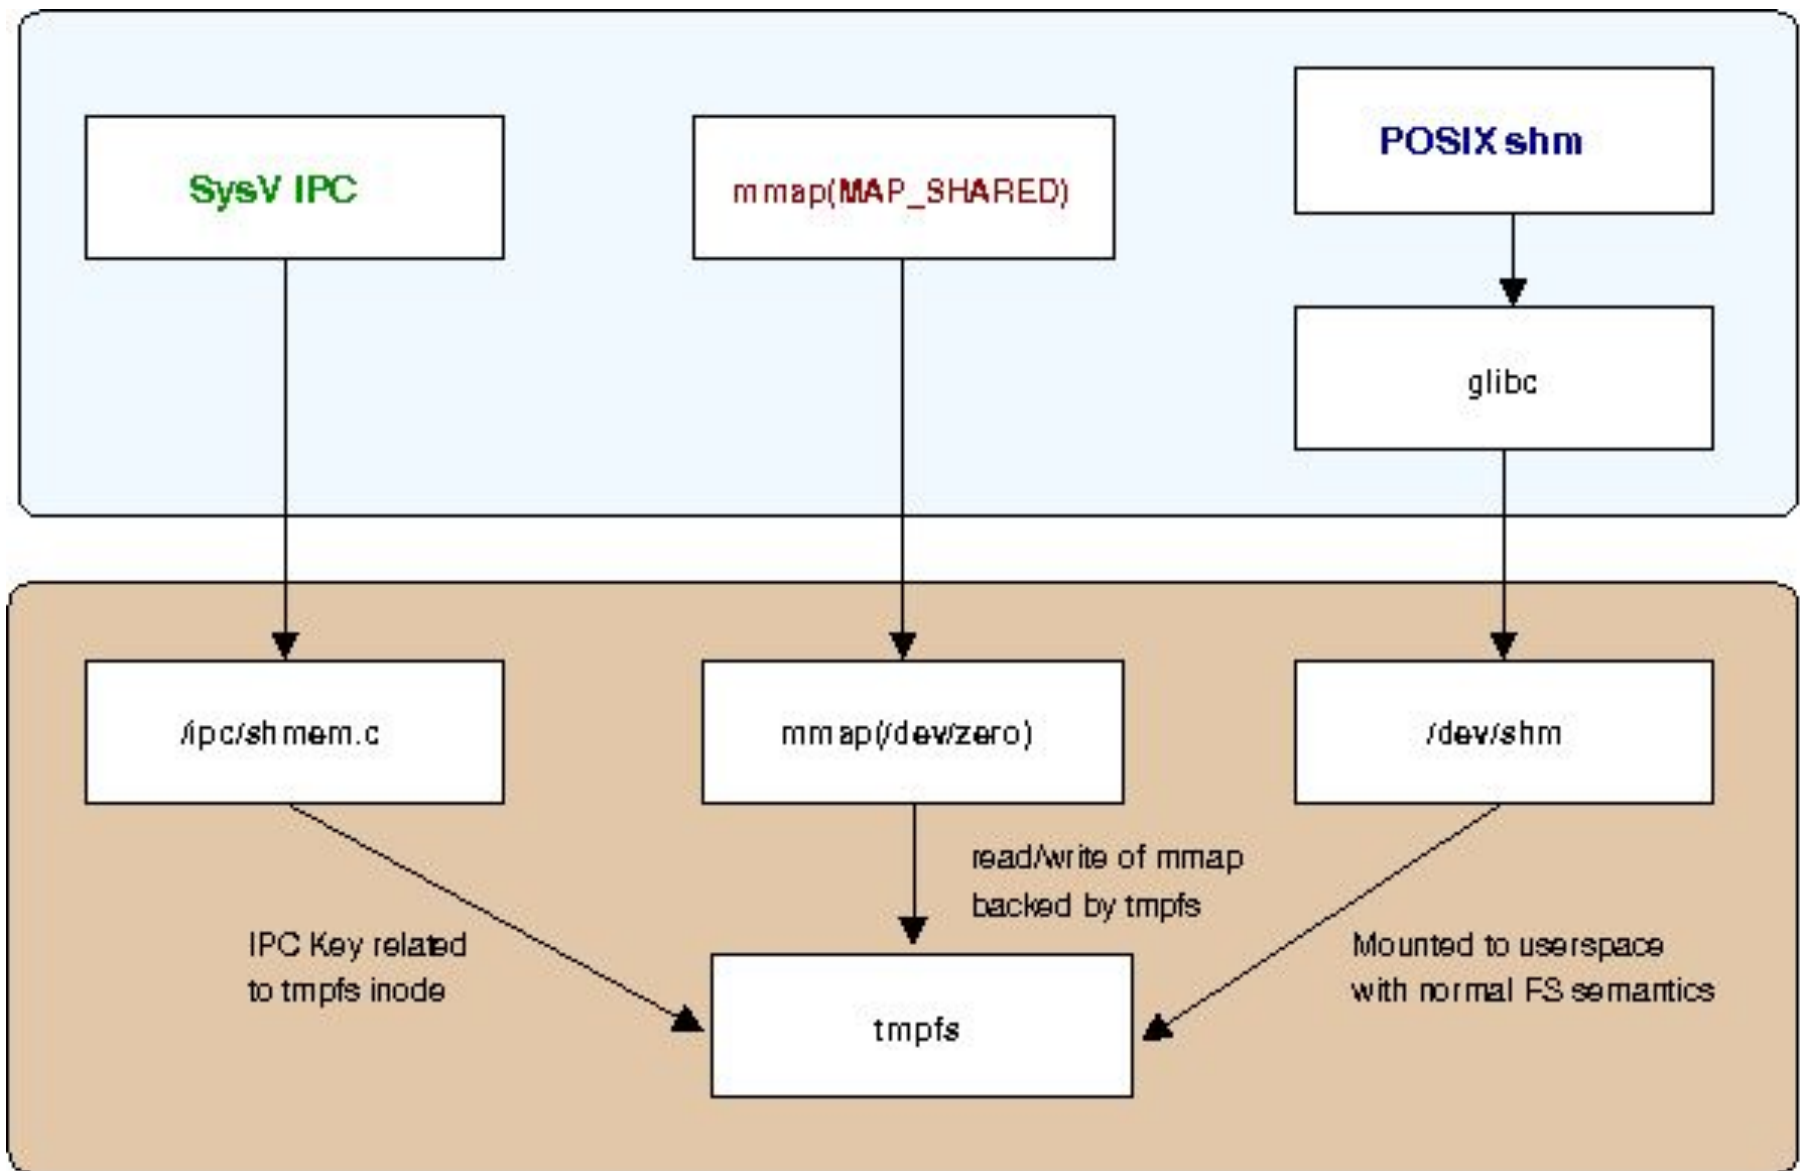
System V	Message queues	Semaphores	Shared memory
Object handle	msqid	semid	shmid
Create/open	msgget()	semget()	shmget() + shmat()
Control/Close	msgctl()	semctl()	shmctl / shmdt()
Perform IPC	msgsnd / msgrcv()	semop()	R/W
POSIX	Message queues	Semaphores	Shared memory
Object handle	mqd_t	sem_t *	fd
Create/open	mq_open()	sem_open()	shm_open() + mmap()
Close	mq_close()	sem_close()	munmap()
Perform IPC	mq_send /mq_receive()	sem_post / sem_wait()	R/W

System V vs. POSIX

POSIX IPC:

Notable Differences

- Uses names, not keys (e.g., `shm_open`)
 - More in line with the UNIX model
- Uses reference counting (e.g., `shm_unlink`)
 - Easier to deallocate resources (and no `ipcs/ipcrm`)
- Provides thread safety
 - Easier to mix multiprocessing and multithreading
- Shared memory is file-oriented
 - Closer to other shared memory interfaces
 - Closer to native `tmpfs` implementation
 - `tmpfs` provides memory-backed filesystem with size limits + swap support



Modern Linux Shmem / tmpfs

References

- [1] Bovet, Daniel P., and Marco Cesati. Understanding the Linux Kernel, 2005.
- [2] Gorman, Mel. Understanding the Linux Virtual Memory Manager, 2004.
- [3] “Linux Cross Reference,” n.d. <https://elixir.bootlin.com/linux/latest/source>.
- [4] “Time Keeping,” n.d. <https://www.kernel.org/doc/Documentation/timers/timekeeping.txt>.
- [5] “CFS Scheduler,” n.d.
<https://www.kernel.org/doc/Documentation/scheduler/sched-design-CFS.txt>.
- [6] “Reducing Scheduling-Clock Ticks,” n.d.
https://www.kernel.org/doc/Documentation/timers/NO_HZ.txt.